

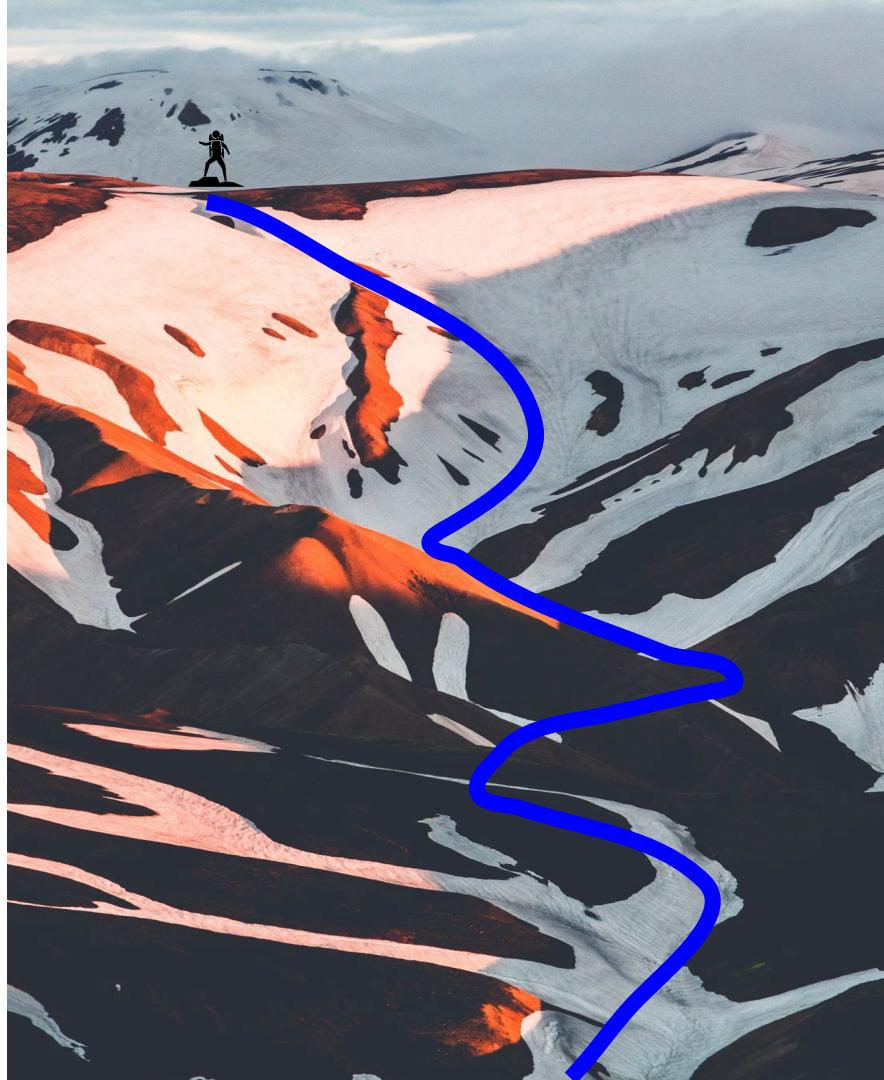
PPGEEC2318

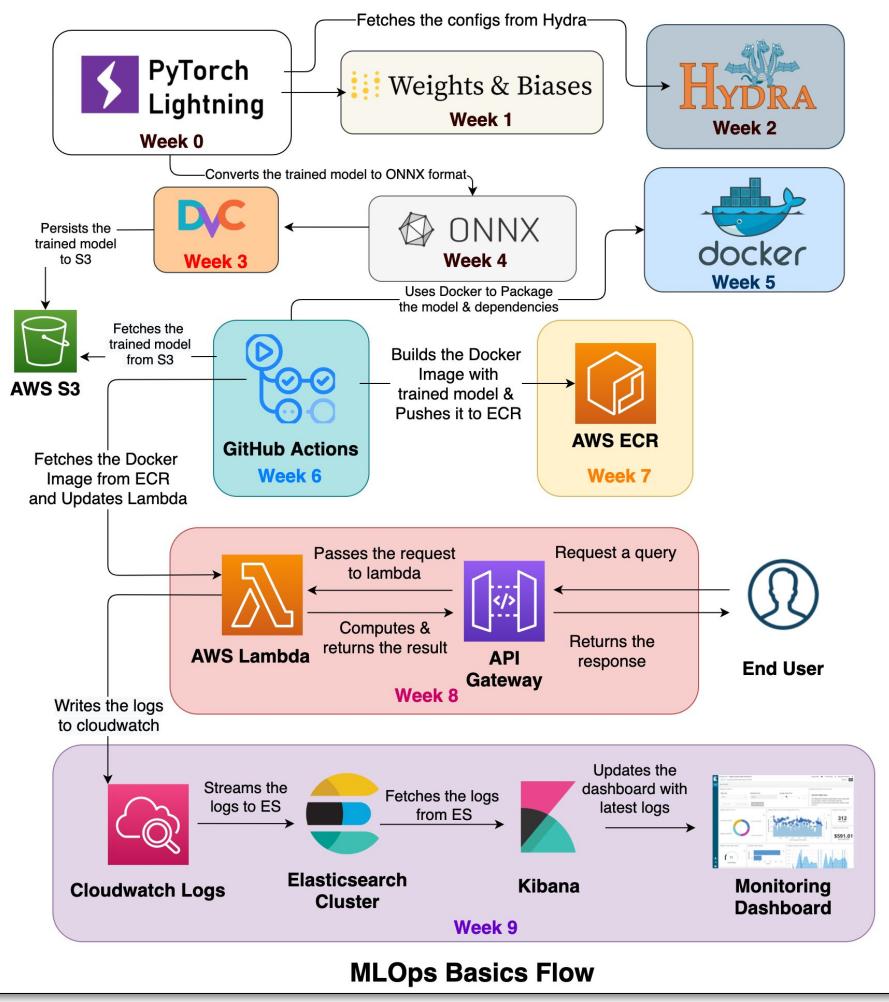
# Machine Learning

Visualizing Gradient Descent  
Regression Problem

Ivanovitch Silva

[ivanovitch.silva@ufrn.br](mailto:ivanovitch.silva@ufrn.br)

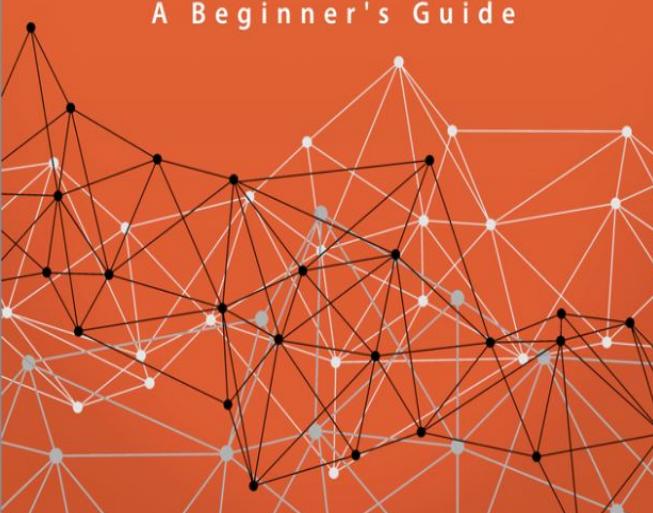




- **Week 0 :** Project Setup - Data acquisition, processing, model declaration, training, and inference.
- **Week 1:** Model Monitoring - Logging, metrics, and visualization using Weights & Biases.
- **Week 2 :** Configurations - Managing ML configurations efficiently with Hydra.
- **Week 3:** Data Version Control - Dataset and model versioning using DVC.
- **Week 4 :** Model Packaging - Converting models to ONNX for cross-platform compatibility.
- **Week 5 :** Model Packaging - Containerizing models and applications with Docker.
- **Week 6 :** CI/CD - Automating ML workflows with GitHub Actions.
- **Week 7 :** Container Registry - Storing and managing container images in AWS ECR.
- **Week 8 :** Serverless Deployment - Deploying ML models using AWS Lambda.
- **Week 9 :** Prediction Monitoring - Setting up Kibana for real-time prediction monitoring.

# Deep Learning with PyTorch Step-by-Step

A Beginner's Guide



Deep Learning with PyTorch Step-by-Step: A Beginner's Guide

Table of Contents

Preface

Acknowledgements

About the Author

- › Frequently Asked Questions (FAQ)
- › Setup Guide
- ▼ Part I: Fundamentals
  - › Chapter 0: Visualizing Gradient Descent
  - › Chapter 1: A Simple Regression Problem

[https://github.com/dvgodoy/  
PyTorchStepByStep](https://github.com/dvgodoy/PyTorchStepByStep)



**Part I: Fundamentals** (gradient descent, training linear and logistic regressions in PyTorch)

**Part II: Computer Vision** (deeper models and activation functions, convolutions, transfer learning, initialization schemes)

**Part III: Sequences** (RNN, GRU, LSTM, seq2seq models, attention, self-attention, transformers)

**Part IV: Natural Language Processing** (tokenization, embeddings, contextual word embeddings, ELMo, BERT, GPT-2)

master

12 Branches 11 Tags

Go to file

Code

 dvgodoy	fixing typo	a6fdc96 · last week	141 Commits
 data_generation	fixing type	last month	
 data_preparation	preview 2.1	4 years ago	
 images	readme	2 years ago	
 model_configuration	revision	3 years ago	
 model_training	revision	3 years ago	
 plots	fixing warnings	last month	
 runs	runs folder	4 years ago	
 stepbystep	revision	3 years ago	
 .gitignore	runs folder	4 years ago	
 Chapter00.ipynb	revision	3 years ago	
 Chapter01.ipynb	revision	2 years ago	
 Chapter02.1.ipynb	revision	3 years ago	
 Chapter02.ipynb	revision	2 years ago	

## About

Official repository of my book: "Deep Learning with PyTorch Step-by-Step: A Beginner's Guide"

 [pytorchstepbystep.com](http://pytorchstepbystep.com)

 python  deep-learning  pytorch  
 pytorch-tutorial  rnn-pytorch  
 cnn-pytorch

 Readme

 MIT license

 Activity

 686 stars

 11 watching

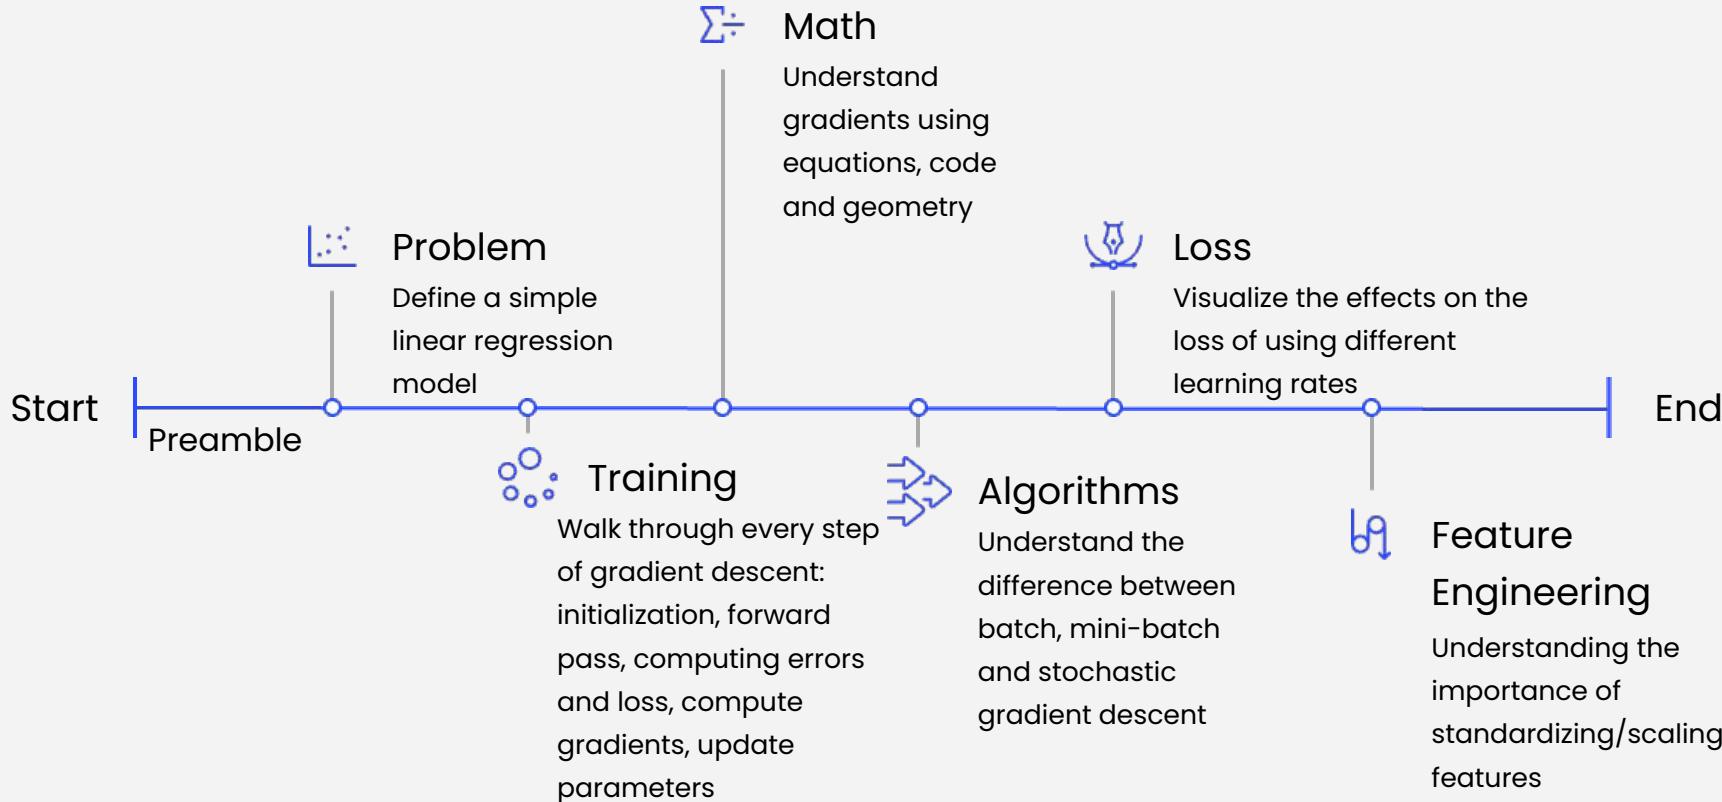
 257 forks

[Report repository](#)

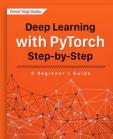
## Releases

 11 tags

## Packages



# Visualizing Gradient Descent



## Preamble

$$y = 1x^2 + 2x + 3$$

$$y = w_1x^2 + w_2x + w_3$$

# Preamble



```
import numpy as np # Import the NumPy library

# Define a 2nd degree polynomial: 1*x^2 + 2*x + 3
polynomial = np.poly1d([1, 2, 3])

# Print the polynomial
print(polynomial)

# Set the number of samples
n = 20

# Generate 'n' random samples from a normal distribution and scale by 5
x = np.random.randn(n, 1) * 5

# Apply the polynomial to the values of x to get corresponding y values
y = polynomial(x)
```

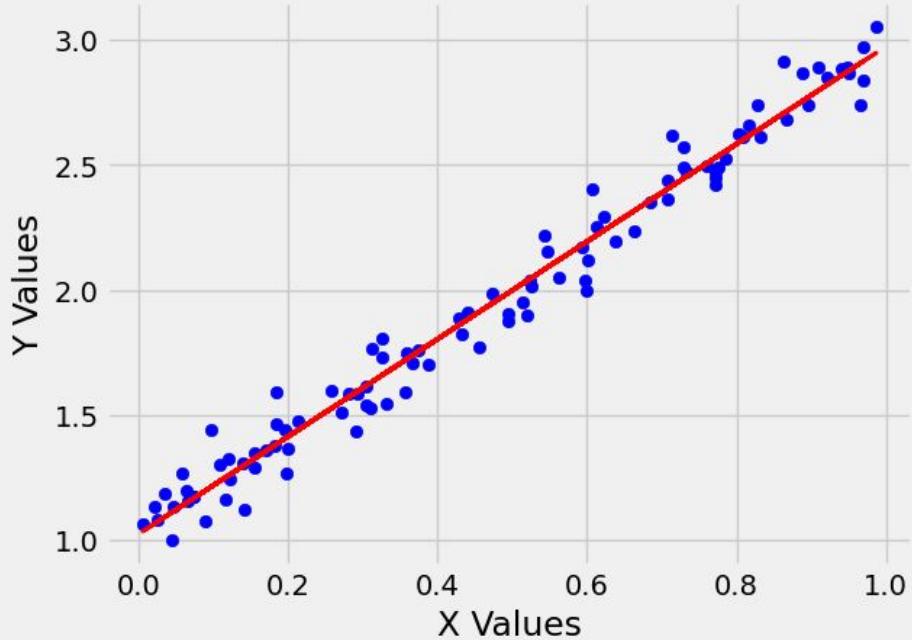
# Preamble

```
# Assume samples X and Y are prepared elsewhere
xx = np.hstack([x*x, x, np.ones_like(x)])
w = torch.randn(3, 1, requires_grad=True) # the 3 coefficients
x = torch.tensor(xx, dtype=torch.float32) # input sample
y = torch.tensor(y, dtype=torch.float32) # output sample
optimizer = torch.optim.NAdam([w], lr=0.01)

# Run optimizer
for _ in range(1000):
    optimizer.zero_grad()
    y_pred = x @ w
    mse = torch.mean(torch.square(y - y_pred))
    mse.backward()
    optimizer.step()
print(w)
tensor([[1.0050],
        [1.9969],
        [2.6778]], requires_grad=True)
```

$$y = w_1 x^2 + w_2 x + w_3$$

Scatter Plot with Linear Regression Line

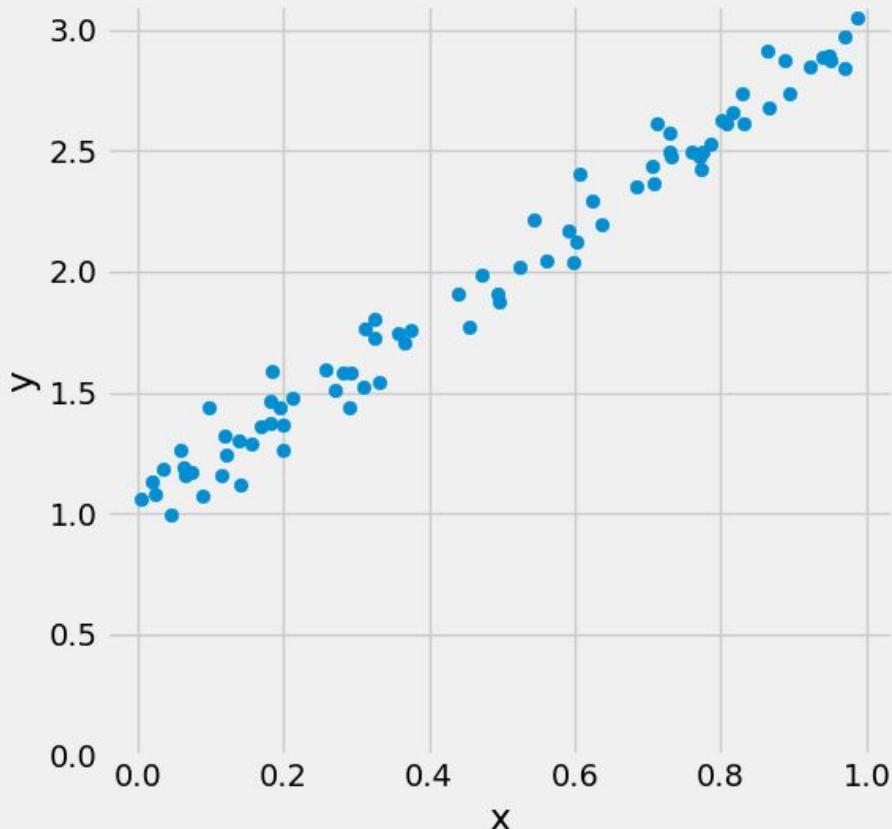


A linear regression model

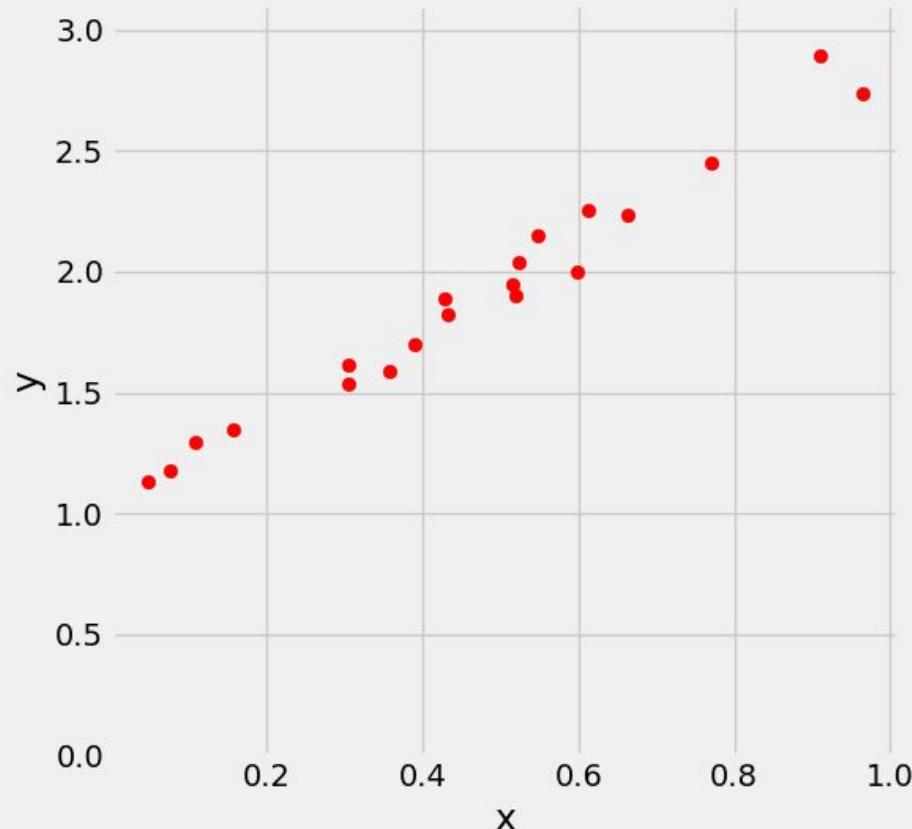
$$y = b + wx + \epsilon$$

# Synthetic Data Generation and Train-Validation-Test Split

Generated Data - Train



Generated Data - Validation





## Step

Initializes parameters

$$y = b + wx$$

```
# Step 0 - Initializes parameters "b" and "w" randomly

# Sets the seed for the random number generator to 42 for reproducibility
np.random.seed(42)

# Generates a single random number from a standard normal distribution for the bias 'b'
b = np.random.randn(1)

# Generates a single random number from a standard normal distribution for the weight 'w'
w = np.random.randn(1)

# Prints the generated values of 'b' and 'w'
print(b, w)

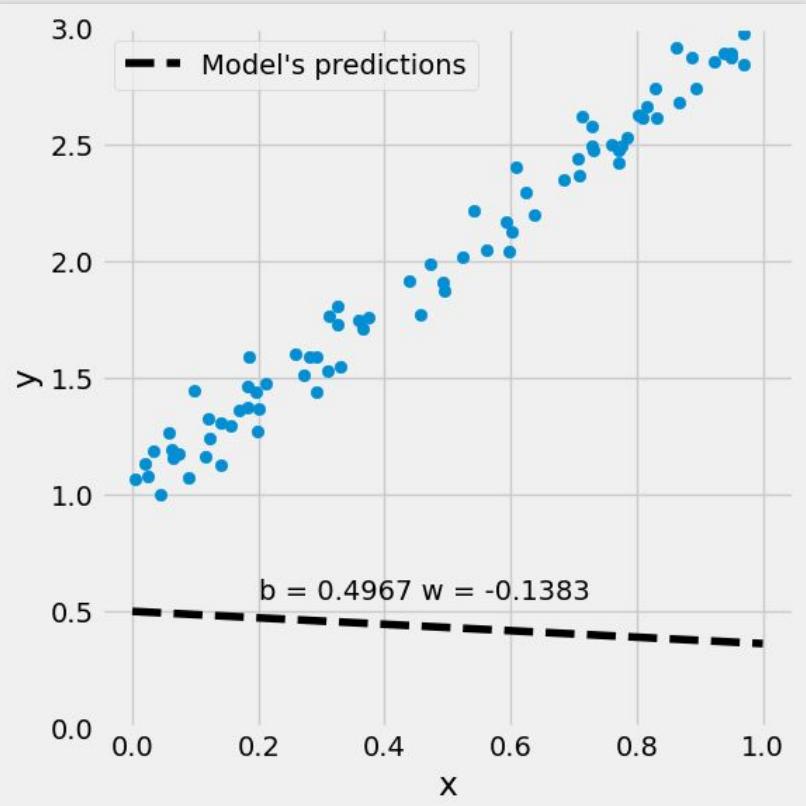
[0.49671415] [-0.1382643]
```

# 01

## Step

Compute predictions

```
# Step 1 - Computes our model's  
predicted output - forward pass  
yhat = b + w * x_train
```



# 02

## Step

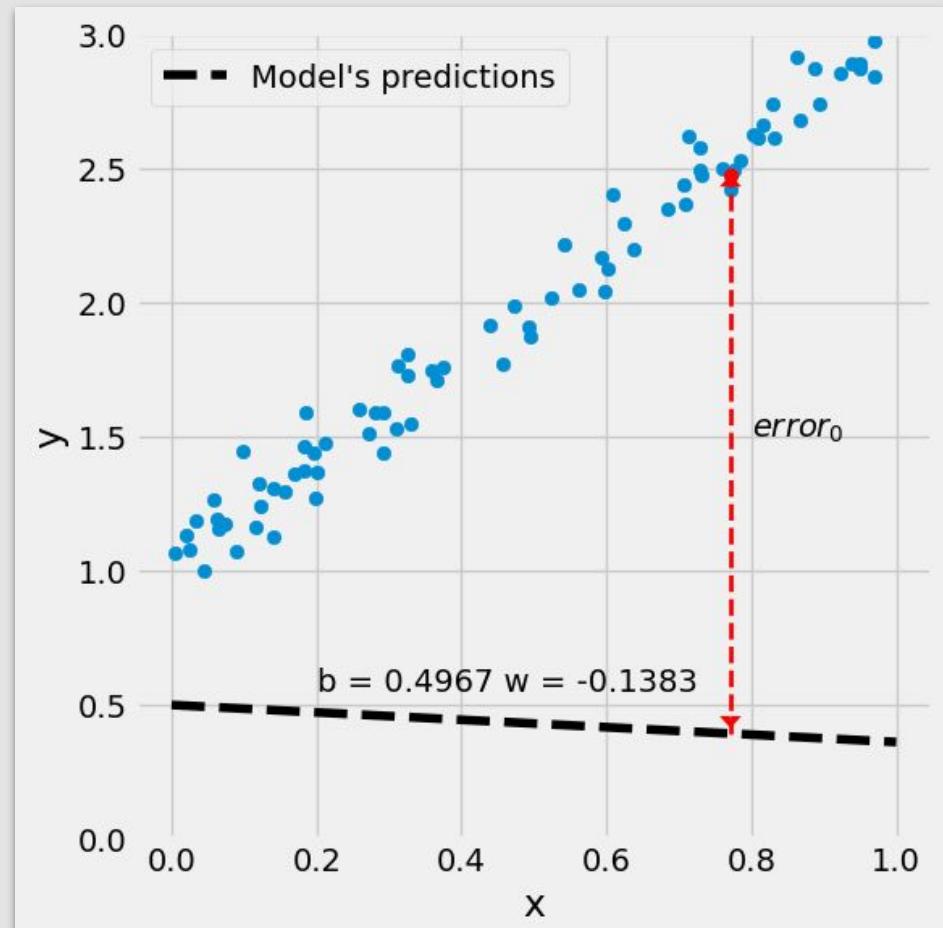
Compute the loss

$$\begin{aligned} \text{MSE} &= \frac{1}{n} \sum_{i=1}^n \text{error}_i^2 \\ &= \frac{1}{n} \sum_{i=1}^n (\hat{y}_i - y_i)^2 \\ &= \frac{1}{n} \sum_{i=1}^n (b + wx_i - y_i)^2 \end{aligned}$$

```
# Step 2 - Computing the loss
# We are using ALL data points, so this is BATCH
gradient
# descent. How wrong is our model? That's the
error!
error = (yhat - y_train)

# It is a regression, so it computes mean squared
error (MSE)
loss = (error ** 2).mean()
print(loss)

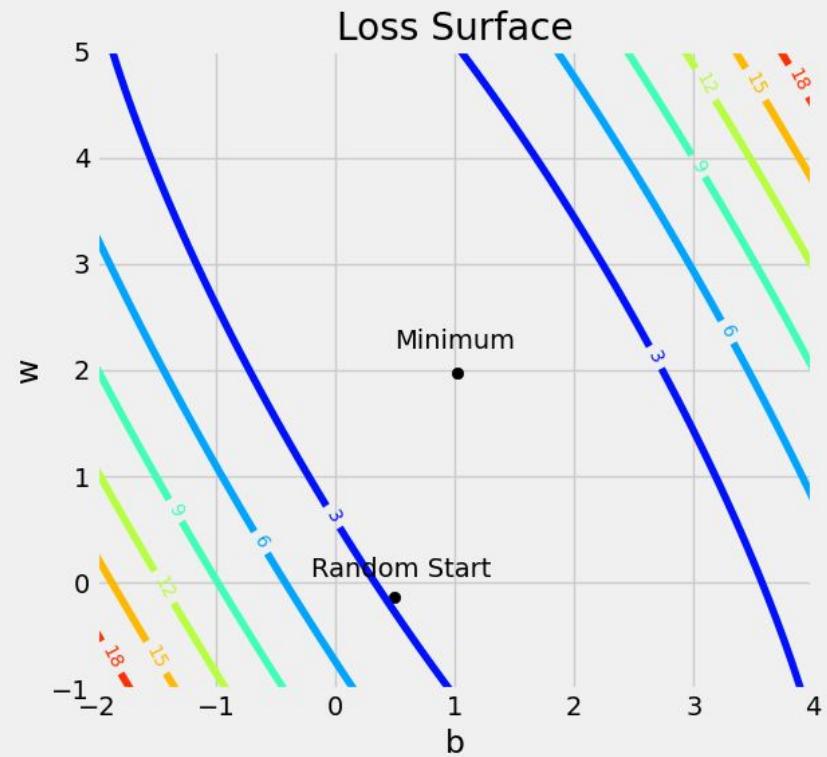
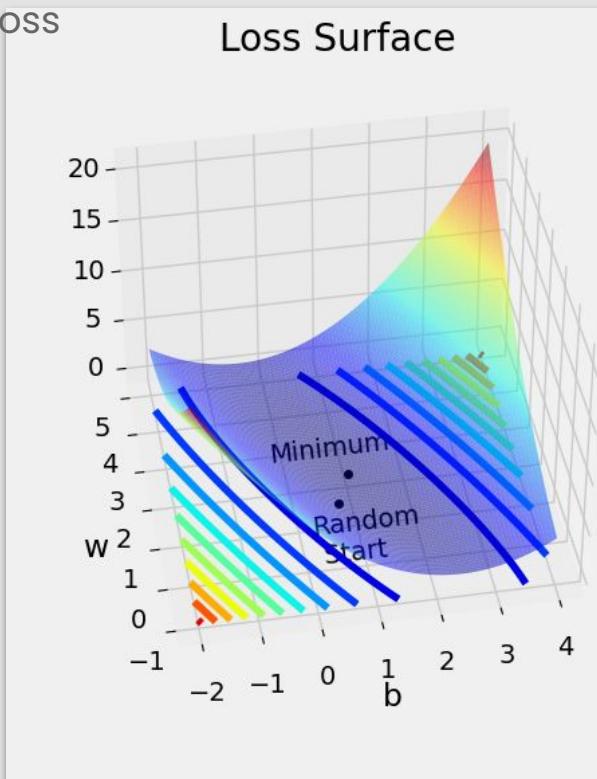
2.7421577700550976
```



# 02

## Step

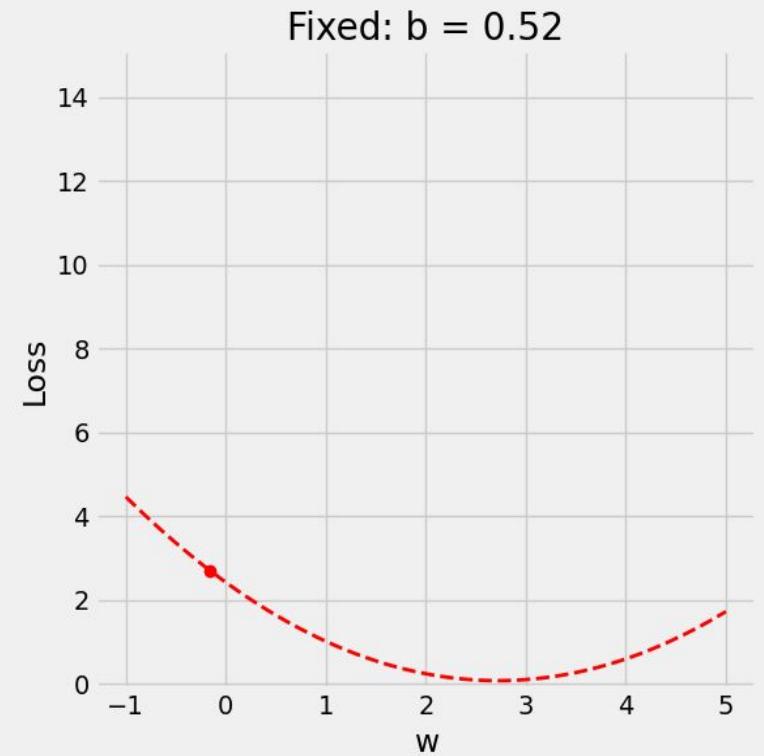
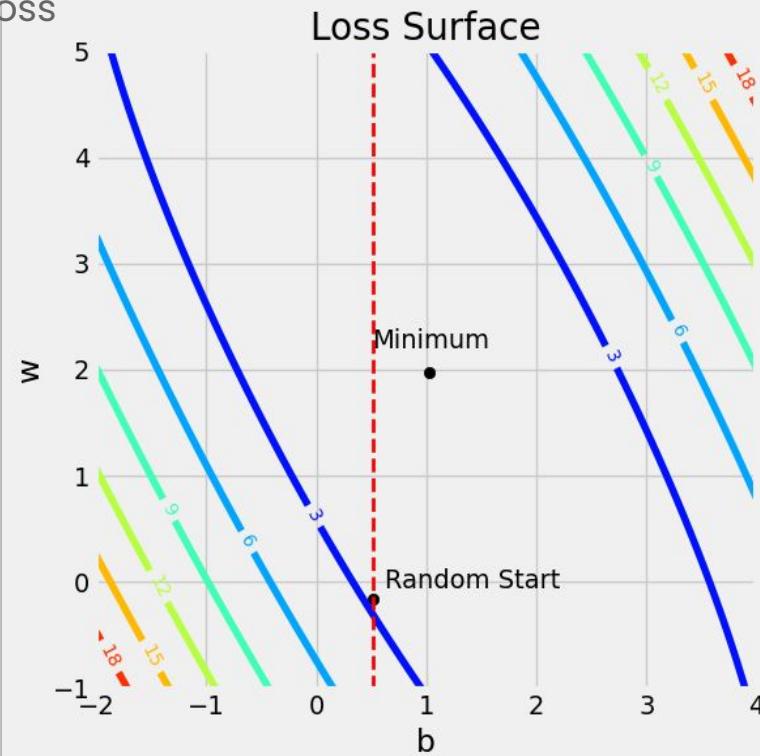
Compute the loss



# 02

## Step

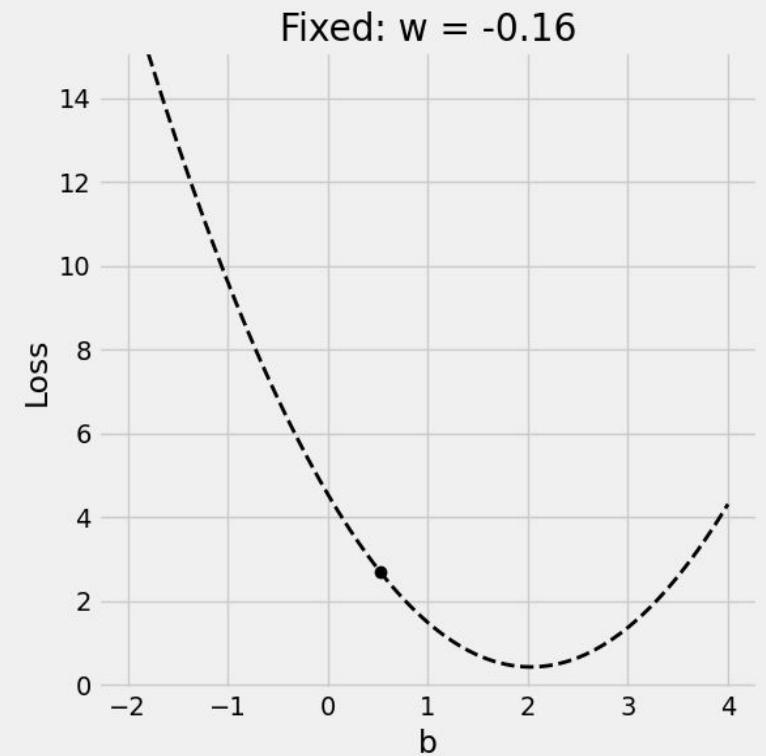
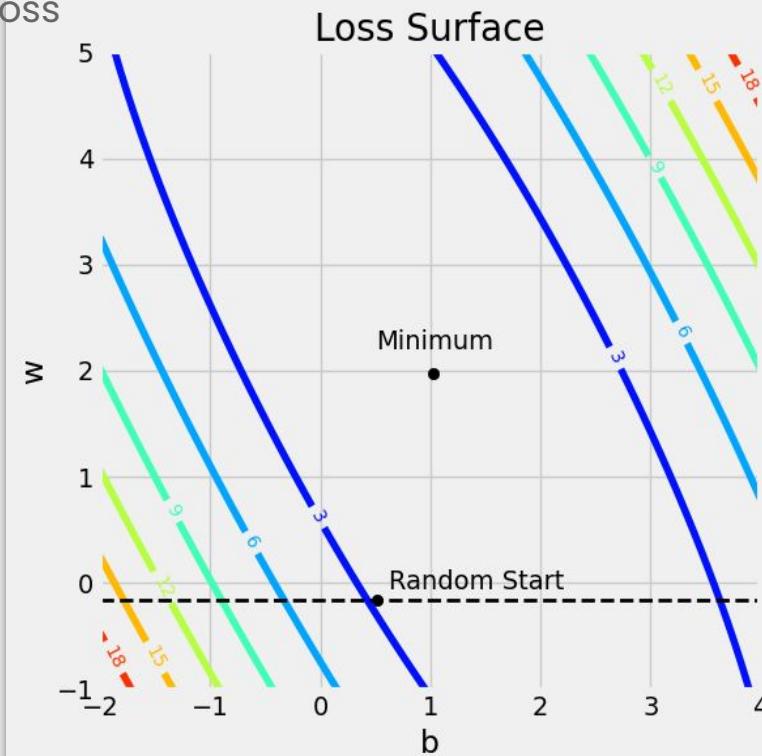
Compute the loss



# 02

## Step

Compute the loss

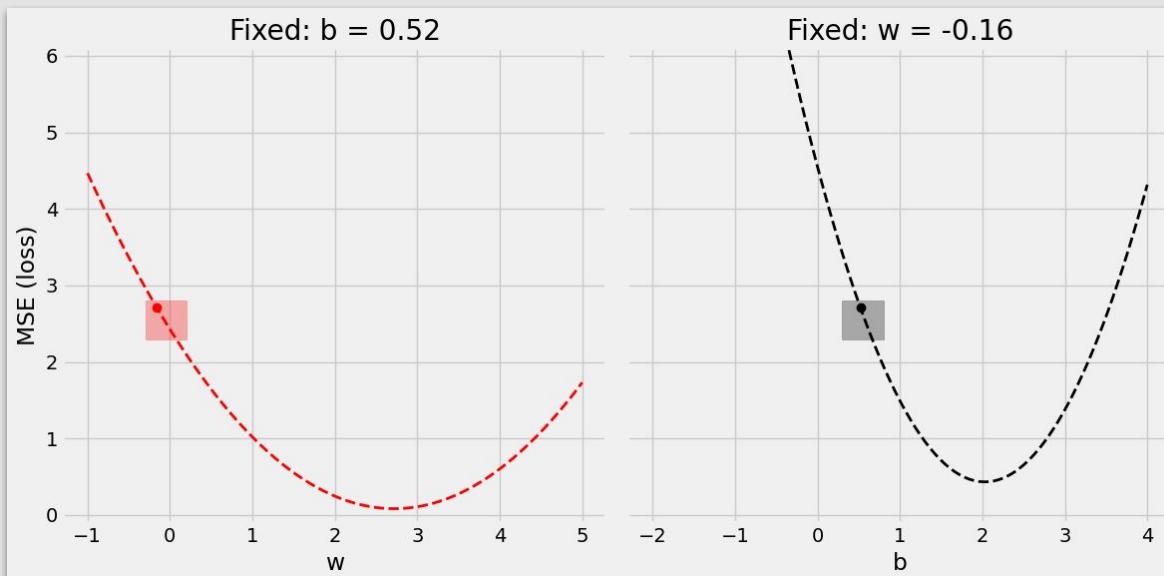


# 03

## Step

Compute the gradients

**Gradient** = how much the loss changes if ONE parameter changes a little bit!



A derivative tells you how much a given quantity changes when you slightly vary some other quantity. In our case, how much does our **MSE loss change** when we **vary each of our two parameters** separately?

# 03

## Step

Compute the gradients

$$\begin{aligned}\text{MSE} &= \frac{1}{n} \sum_{i=1}^n \text{error}_i^2 \\ &= \frac{1}{n} \sum_{i=1}^n (\hat{y}_i - y_i)^2 \\ &= \frac{1}{n} \sum_{i=1}^n (b + wx_i - y_i)^2\end{aligned}$$

$$\begin{aligned}\frac{\partial \text{MSE}}{\partial b} &= \frac{\partial \text{MSE}}{\partial \hat{y}_i} \frac{\partial \hat{y}_i}{\partial b} = \frac{1}{n} \sum_{i=1}^n 2(b + wx_i - y_i) \\ &= 2 \frac{1}{n} \sum_{i=1}^n (\hat{y}_i - y_i)\end{aligned}$$

$$\begin{aligned}\frac{\partial \text{MSE}}{\partial w} &= \frac{\partial \text{MSE}}{\partial \hat{y}_i} \frac{\partial \hat{y}_i}{\partial w} = \frac{1}{n} \sum_{i=1}^n 2(b + wx_i - y_i)x_i \\ &= 2 \frac{1}{n} \sum_{i=1}^n x_i(\hat{y}_i - y_i)\end{aligned}$$

```
# Step 3 - Computes gradients for both "b" and "w" parameters
b_grad = 2 * error.mean()
w_grad = 2 * (x_train * error).mean()
print(b_grad, w_grad)
```

```
-3.044811379650508 -1.8337537171510832
```

# 04

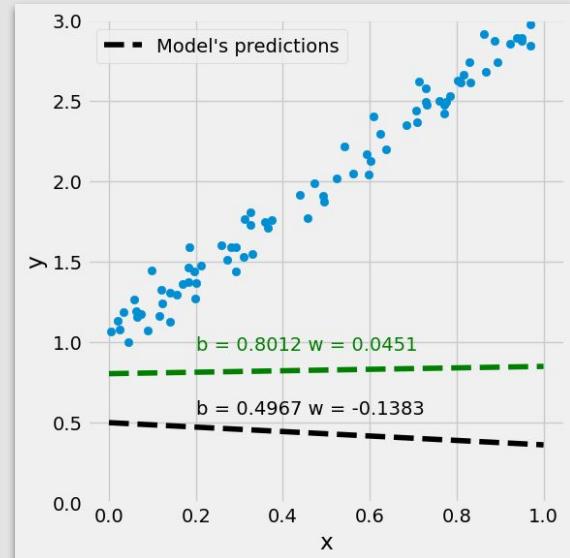
## Step

Update the parameters

```
# Sets learning rate - this is "eta" ~ the "n" like Greek letter
lr = 0.1
print(b, w)

# Step 4 - Updates parameters using gradients and the
# learning rate
b = b - lr * b_grad
w = w - lr * w_grad

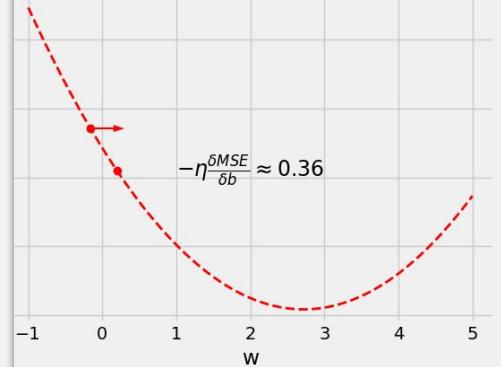
print(b, w)
[0.49671415] [-0.1382643]
[0.80119529] [0.04511107]
```



$$b = b - \eta \frac{\partial \text{MSE}}{\partial b}$$
$$w = w - \eta \frac{\partial \text{MSE}}{\partial w}$$

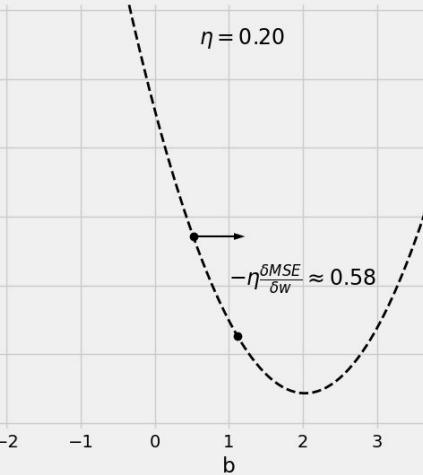
Fixed:  $b = 0.52$

$\eta = 0.20$



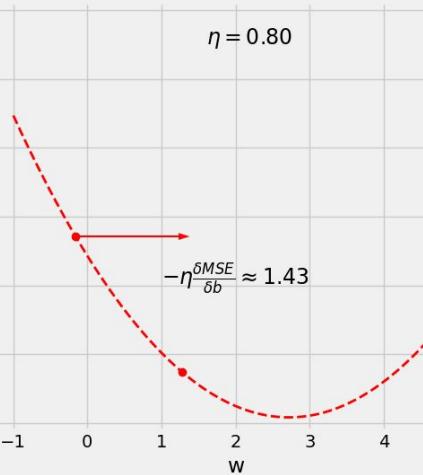
Fixed:  $w = -0.16$

$\eta = 0.20$



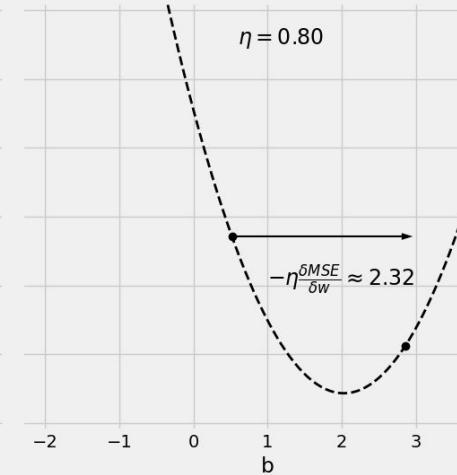
Fixed:  $b = 0.52$

$\eta = 0.80$



Fixed:  $w = -0.16$

$\eta = 0.80$

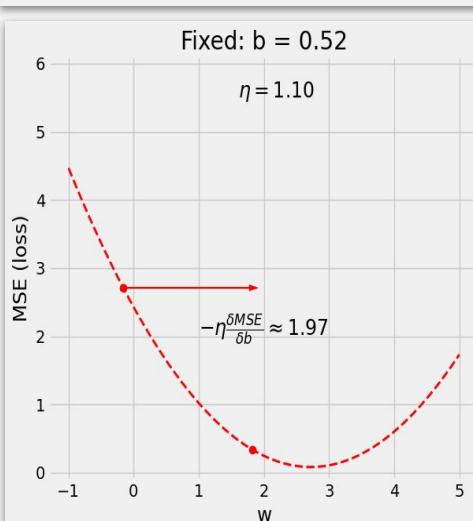


# Learning Rate

Fixed:  $b = 0.52$

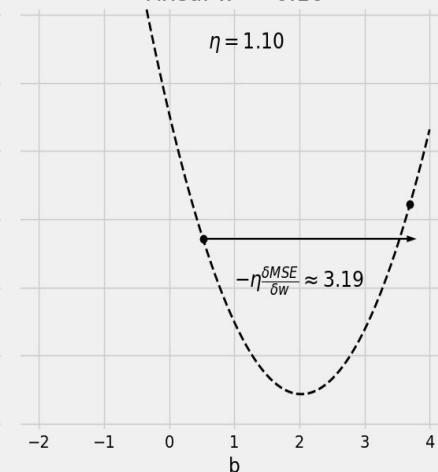
$\eta = 1.10$

MSE (loss)

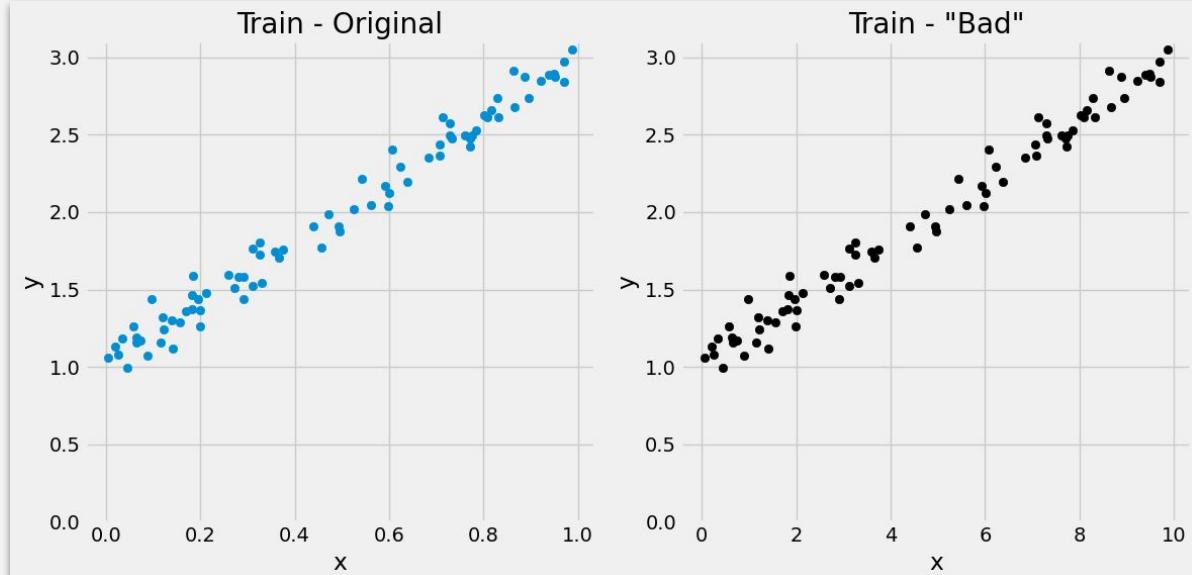


Fixed:  $w = -0.16$

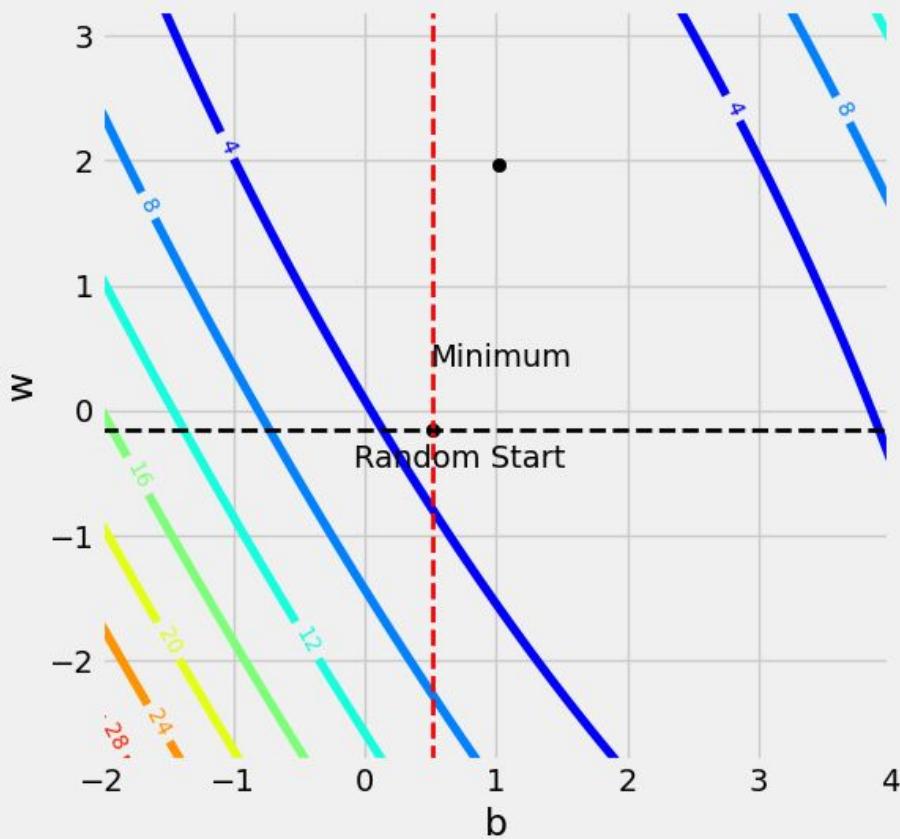
$\eta = 1.10$



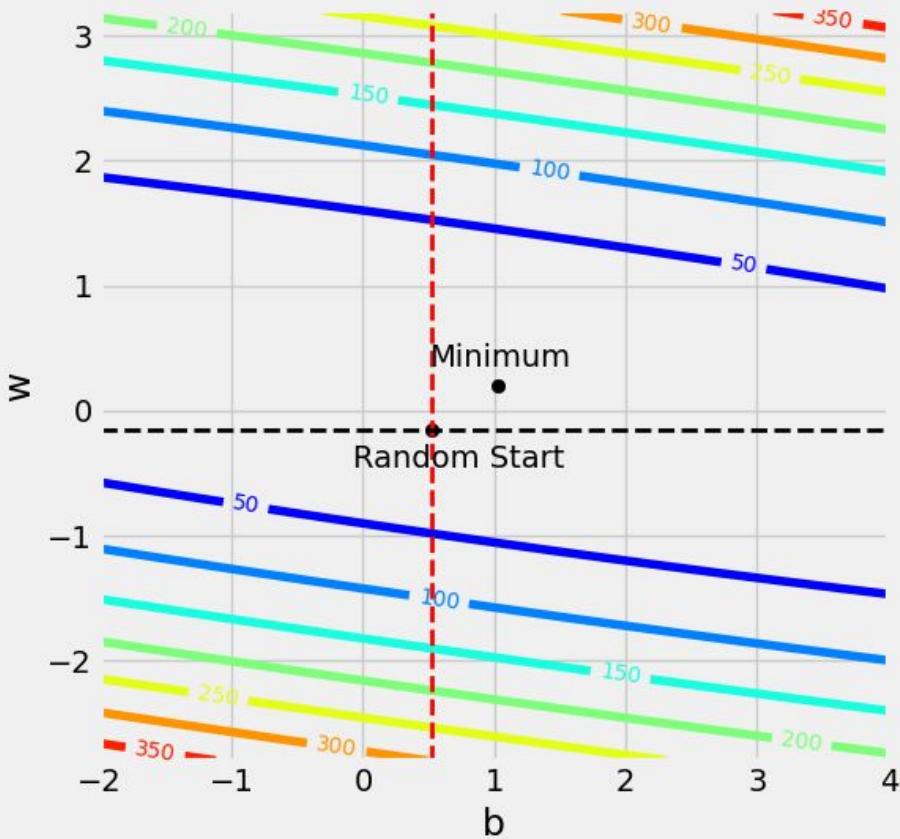
What is the influence of the X scale on the results?



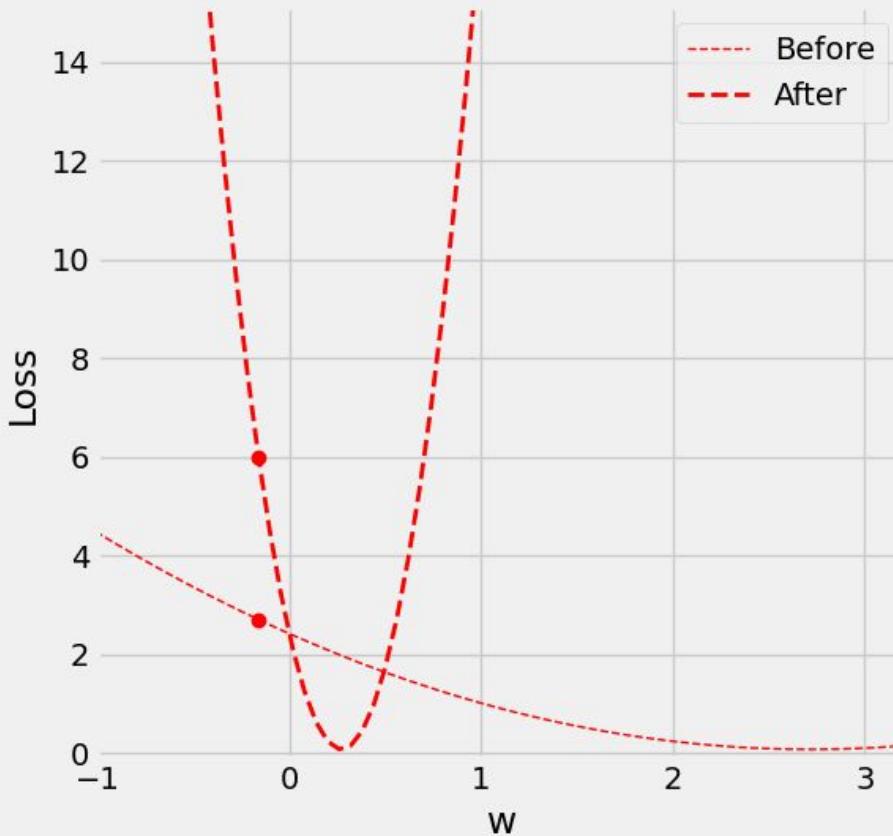
### Loss Surface - Before



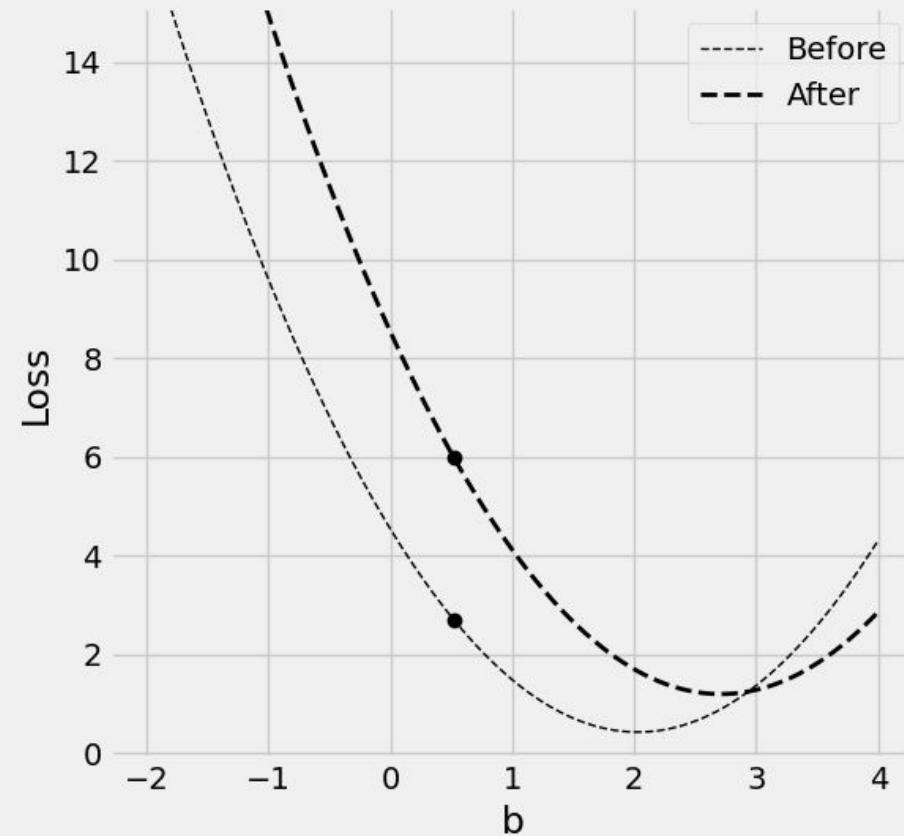
### Loss Surface - After



Fixed:  $b = 0.52$



Fixed:  $w = -0.16$



The red curve got much steeper (larger gradient), and thus we must use a lower learning rate to safely descend along it.

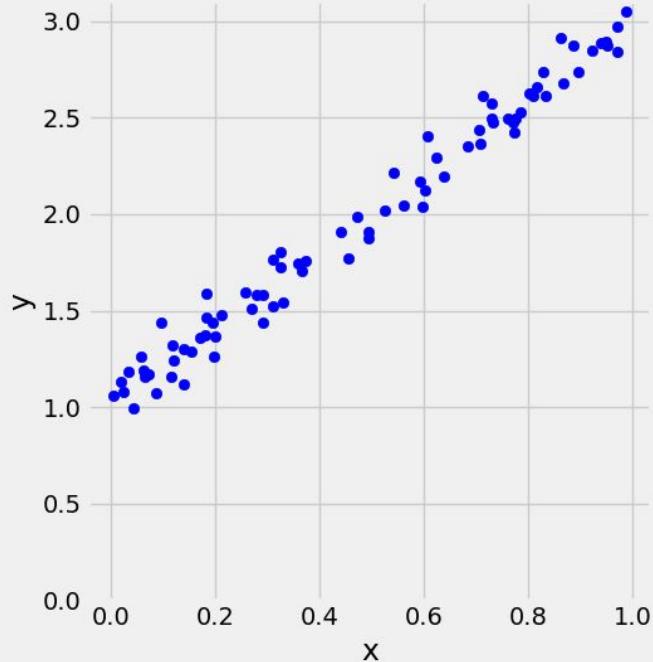
How can  
we fix it?

$$\bar{X} = \frac{1}{N} \sum_{i=1}^N x_i$$

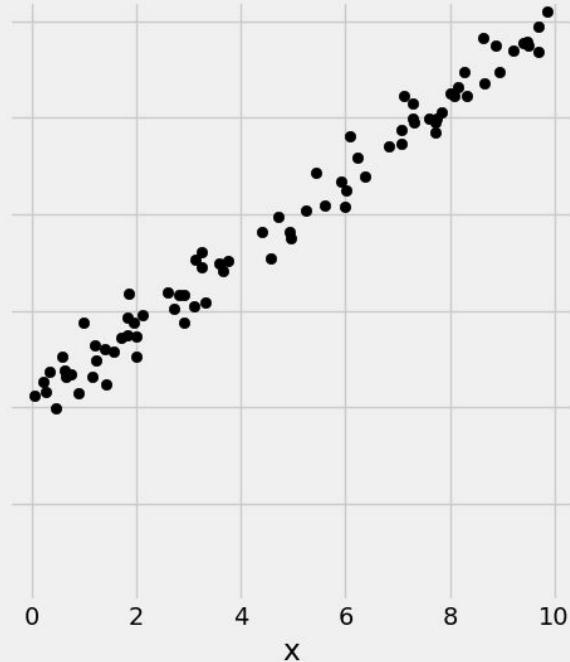
$$\sigma(X) = \sqrt{\frac{1}{N} \sum_{i=1}^N (x_i - \bar{X})^2}$$

$$\text{scaled } x_i = \frac{x_i - \bar{X}}{\sigma(X)}$$

Train - Original

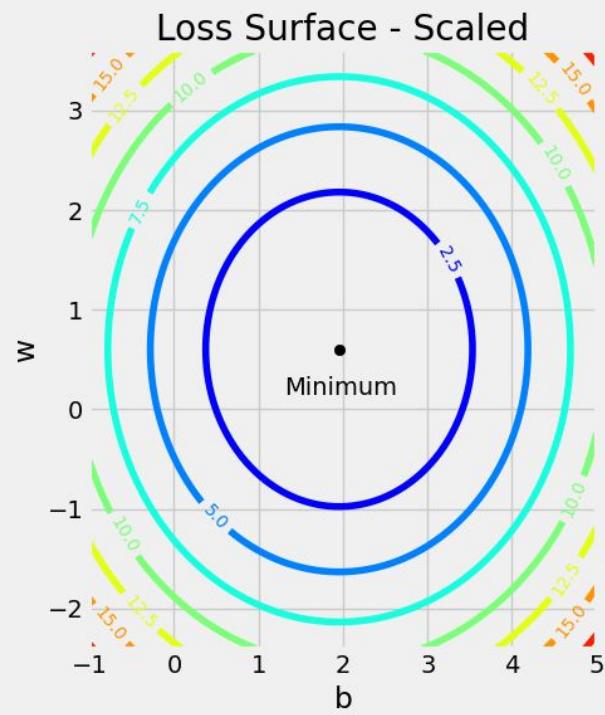
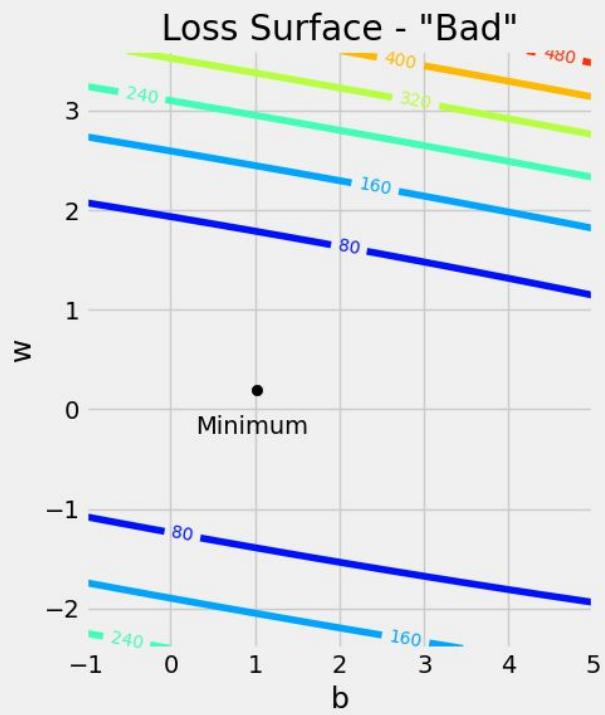
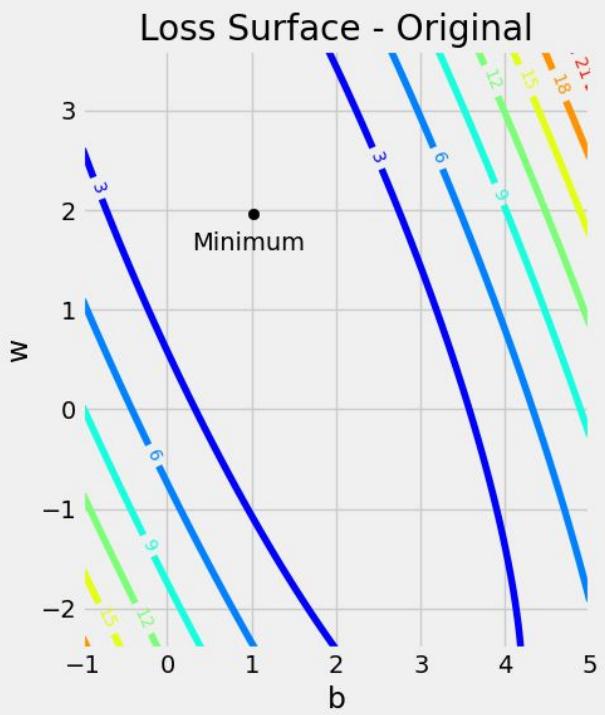


Train - "Bad"



Train - Scaled



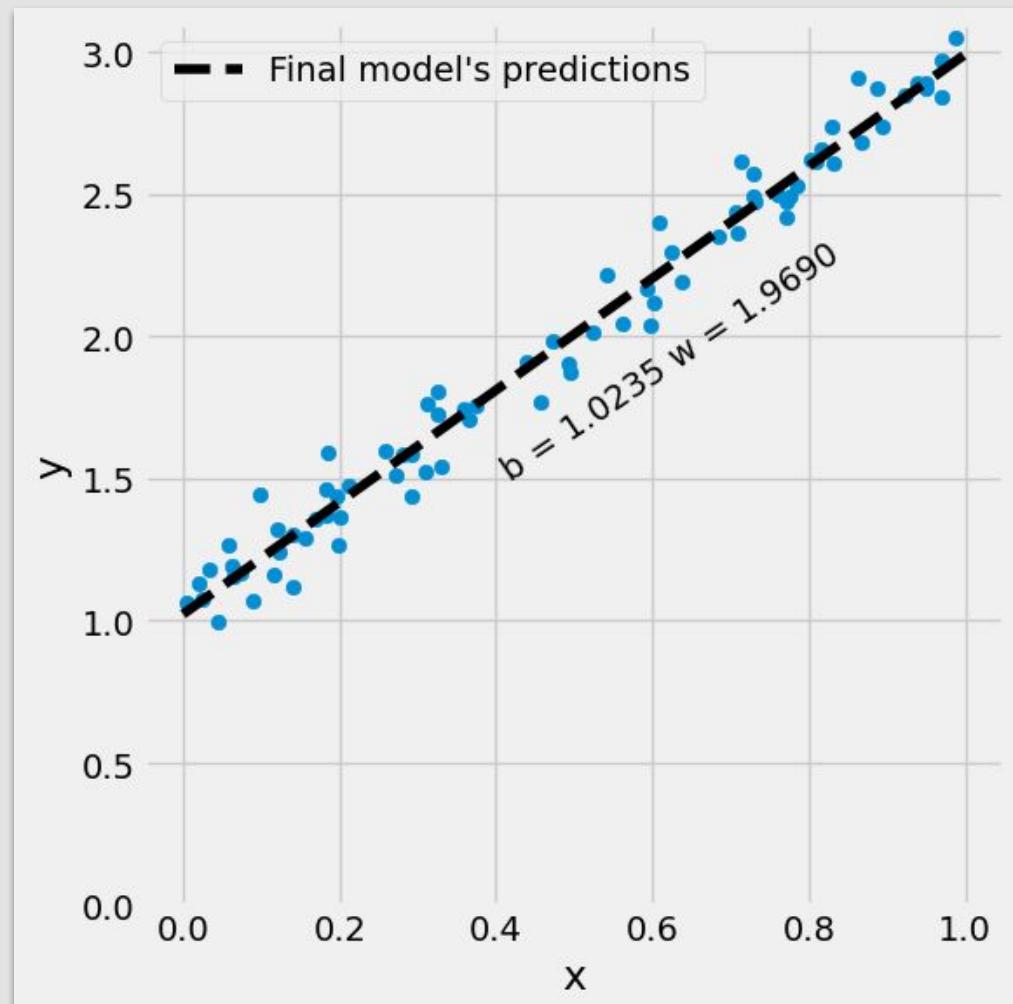


# 05

## Step

Iterate and Converge

Now we use the updated parameters to go back to Step 1 and restart the process.



# Definition of Epoch

An epoch is complete whenever every point in the training set( $N$ ) has already been used in all steps:  
forward pass, computing loss, computing gradients, and updating parameters.

- Batch gradient descent

*This is trivial, as it uses all points for computing the loss  
– one epoch is the same as one update.*

- Stochastic gradient descent

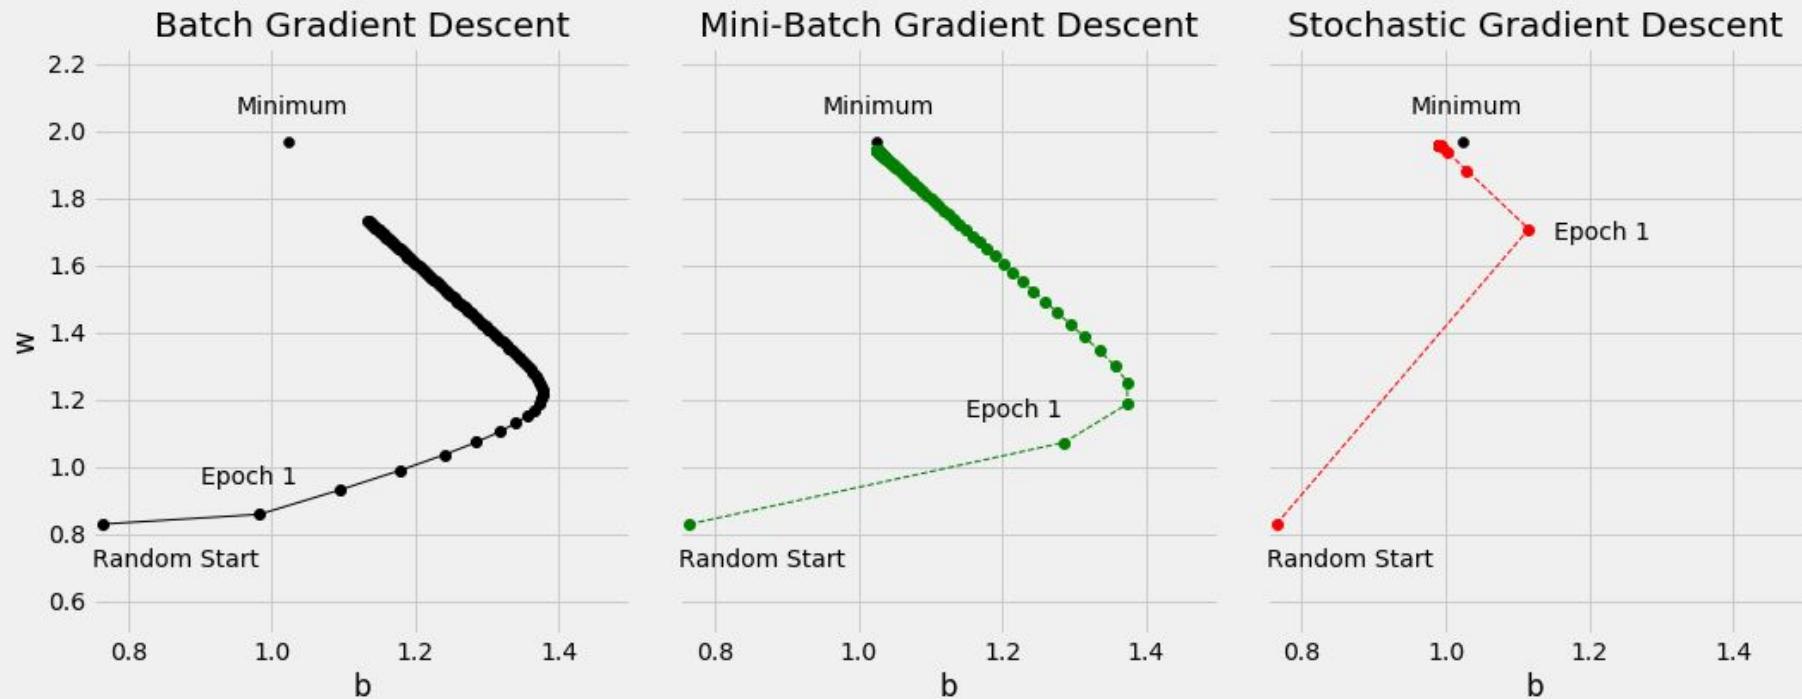
*One epoch means  $N$  updates, since every individual data point is used to perform an update.*

- Mini-batch gradient descent

*One epoch has  $N/n$  updates since a mini-batch of  $n$  data points is used to perform an update.*



100 epochs using either 80 data points (batch), 16 data points (mini-batch), or a single data point (stochastic) for computing the loss, as shown in the figure below.



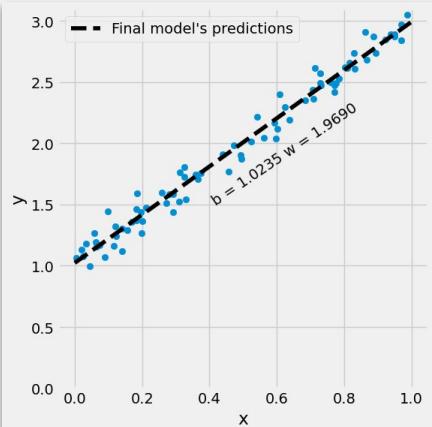
```

# Step 0 - Initializes parameters
# "b" and "w" randomly
np.random.seed(42)
b = np.random.randn(1)
w = np.random.randn(1)

print(b, w)

# Sets learning rate - this is
# "eta" ~ the "n"-like Greek letter
lr = 0.1
# Defines number of epochs
n_epochs = 1000

```



```

for epoch in range(n_epochs):
    # Step 1 - Computes model's predicted output - forward pass
    yhat = b + w * x_train

    # Step 2 - Computes the loss
    # We are using ALL data points, so this is BATCH gradient
    # descent. How wrong is our model? That's the error!
    error = (yhat - y_train)
    # It is a regression, so it computes mean squared error (MSE)
    loss = (error ** 2).mean()

    # Step 3 - Computes gradients for both "b" and "w" parameters
    b_grad = 2 * error.mean()
    w_grad = 2 * (x_train * error).mean()

    # Step 4 - Updates parameters using gradients and
    # the learning rate
    b = b - lr * b_grad
    w = w - lr * w_grad

print(b, w)
[0.49671415] [-0.1382643]
[1.02354094] [1.96896411]

```

**Tom Yeh** (He/Him) • Following

Associate Professor of Computer Science at University of Colorado...

2d • 🌐

...

### 3. Linear Layer: 25 Exercises 🎉

~ AI by Hand ↗ Workbook ~

#### -- Previous Workbooks --

1: Dot Product

<https://lnkd.in/g8TNrpfN>

2: Matrix Multiplication

<https://lnkd.in/g6DnVySK>

#### -- Workbook --

Each workbook is comprised of 25 exercises.

Every exercise features a "missing" value, highlighted in a soft shade of red.

Can you calculate this missing value by hand? 🤝

For each exercise, the solution can be found in the bottom-right corner of the page. Do your best to avoid looking ahead. ☺

#### -- Share --

#linear #aibyhand #ai #neuralnetwork

[Repost 🌱] Please help me share this workbook with more people!

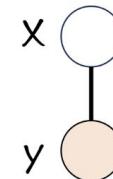
Linear Layer: AI by Hand ↗ Workbook • 25 pages

Linear Layer

### Exercise 1

$$\begin{matrix} X \\ 4 \\ \omega & b & 1 \\ 3 & 0 & y \end{matrix}$$

$$y = [\omega|b] \cdot [X|1] = \omega X + b$$

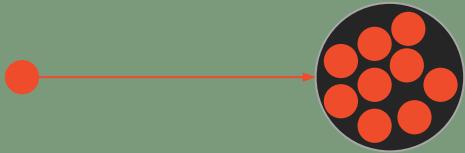


**This is pure gold!!!**

12

AI by Hand ↗ © 2024 Tom Yeh

# *Time to TORCH it!*





# Pytorch Kickoff

Pytorch is more pythonic? TensorFlow excels in large-scale and production environments?

## Fundamentals

Tensors, Loading Data, Devices and CUDA, Creating Parameters

## Autograd

Backward, Update Parameters, Dynamic Computational Graph

## Putting It All Together

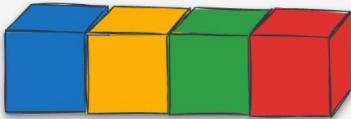
Optimizer, Loss, Models, Training

```
import torch
import torch.optim as optim
import torch.nn as nn
from torchviz import make_dot
```

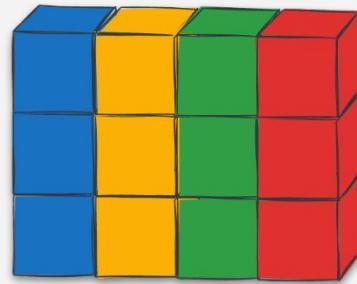
In the world of PyTorch, every data point, image, or signal takes shape as a tensor



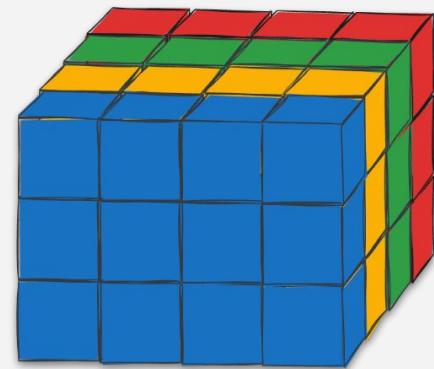
Scalar



Vector



Matrix



Tensor

```
scalar = torch.tensor(3.14159)
vector = torch.tensor([1, 2, 3])
matrix = torch.ones((2, 3), dtype=torch.float)
tensor = torch.randn((2, 3, 4), dtype=torch.float)

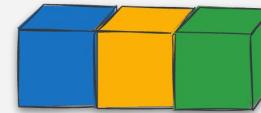
print(scalar)
print(vector)
print(matrix)
print(tensor)

tensor(3.1416)
tensor([1, 2, 3])
tensor([[1., 1., 1.],
       [1., 1., 1.]])
tensor([[[ 0.9755, -0.8582, -1.5437,  0.8710],
        [ 2.2040,  0.2876,  1.4933, -1.6427],
        [ 0.6846,  0.2353, -0.9709, -1.3681]],

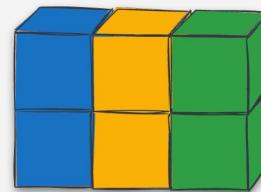
       [[ 4.4695,  0.0782, -0.6045,  1.0868],
        [ 1.1322, -0.4760,  0.6845, -0.4704],
        [ 1.5749, -0.3304, -0.7718,  0.6357]]])
```



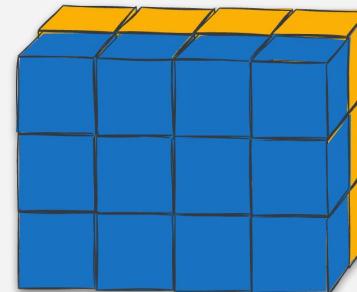
Scalar



Vector



Matrix



Tensor

```
device = 'cuda' if torch.cuda.is_available() else 'cpu'

# Our data was in Numpy arrays, but we need to transform them
# into PyTorch's Tensors and then we send them to the
# chosen device
x_train_tensor = torch.as_tensor(x_train).float().to(device)
y_train_tensor = torch.as_tensor(y_train).float().to(device)
```



Loading Data  
& Devices

```
# We can specify the device at the moment of creation
# RECOMMENDED!

# Step 0 - Initializes parameters "b" and "w" randomly
torch.manual_seed(42)
b = torch.randn(1, requires_grad=True, \
                dtype=torch.float, device=device)
w = torch.randn(1, requires_grad=True, \
                dtype=torch.float, device=device)
print(b, w)

tensor([0.1940], device='cuda:0', requires_grad=True)
tensor([0.1391], device='cuda:0', requires_grad=True)
```



## Best-Practices

```

# Step 1 - Computes our model's predicted output - forward pass
yhat = b + w * x_train_tensor

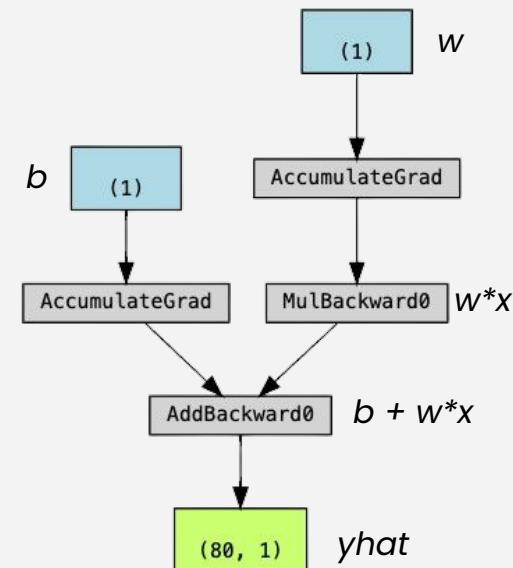
# Step 2 - Computes the loss
# We are using ALL data points, so this is BATCH gradient
descent
# How wrong is our model? That's the error!
error = (yhat - y_train_tensor)
# It is a regression, so it computes mean squared error (MSE)
loss = (error ** 2).mean()

# Step 3 - Computes gradients for both "b" and "w" parameters
# No more manual computation of gradients!
# b_grad = 2 * error.mean()
# w_grad = 2 * (x_tensor * error).mean()
loss.backward()

```

# Autograd

## backward



Dynamic Computational Graph

# Autograd

## Updating Parameters

```
# Defines number of epochs
n_epochs = 1000

for epoch in range(n_epochs):
    # Step 1 - Computes model's predicted output - forward pass
    yhat = b + w * x_train_tensor

    # Step 2 - Computes the loss
    # We are using ALL data points, so this is BATCH gradient
    # descent. How wrong is our model? That's the error!
    error = (yhat - y_train_tensor)
    # It is a regression, so it computes mean squared error (MSE)
    loss = (error ** 2).mean()

    # Step 3 - Computes gradients for both "b" and "w" parameters
    # No more manual computation of gradients!
    # b_grad = 2 * error.mean()
    # w_grad = 2 * (x_tensor * error).mean()
    # We just tell PyTorch to work its way BACKWARDS
    # from the specified loss!
    loss.backward()
```

```
# Step 4 - Updates parameters using gradients and
# the learning rate.
# We need to use NO_GRAD to keep the update out of
# the gradient computation. Why is that? It boils
# down to the DYNAMIC GRAPH that PyTorch uses...
with torch.no_grad():
    b -= lr * b.grad
    w -= lr * w.grad

    # PyTorch is "clingy" to its computed gradients,
    # we need to tell it to let it go...
    b.grad.zero_()
    w.grad.zero_()
print(b, w)

tensor([1.0235], device='cuda:0', requires_grad=True)
tensor([1.9690], device='cuda:0', requires_grad=True)
```

# Optimizer and Loss

```
# Sets learning rate - this is "eta" ~ the "n"-like
# Greek letter
lr = 0.1

# Step 0 - Initializes parameters "b" and "w"
# randomly
torch.manual_seed(42)
b = torch.randn(1, requires_grad=True, \
               dtype=torch.float, device=device)
w = torch.randn(1, requires_grad=True, \
               dtype=torch.float, device=device)

# Defines a SGD optimizer to update the parameters
optimizer = optim.SGD([b, w], lr=lr)

# Defines a MSE loss function
loss_fn = nn.MSELoss(reduction='mean')

# Defines number of epochs
n_epochs = 1000
```

```
for epoch in range(n_epochs):
    # Step 1 - Computes model's predicted output - forward pass
    yhat = b + w * x_train_tensor

    # Step 2 - Computes the loss
    # No more manual loss!
    # error = (yhat - y_train_tensor)
    # loss = (error ** 2).mean()
    loss = loss_fn(yhat, y_train_tensor)

    # Step 3 - Computes gradients for both "b" and "w"
    loss.backward()

    # Step 4 - Updates parameters using gradients and
    # the learning rate
    optimizer.step()
    optimizer.zero_grad()

print(b, w)
tensor([1.0235], device='cuda:0', requires_grad=True)
tensor([1.9690], device='cuda:0', requires_grad=True)
```

```
class ManualLinearRegression(nn.Module):
    def __init__(self):
        super().__init__()
        # To make "b" and "w" real parameters of the model,
        # we need to wrap them with nn.Parameter
        self.b = nn.Parameter(torch.randn(1,
                                         requires_grad=True,
                                         dtype=torch.float))
        self.w = nn.Parameter(torch.randn(1,
                                         requires_grad=True,
                                         dtype=torch.float))

    def forward(self, x):
        # Computes the outputs / predictions
        return self.b + self.w * x
```

# Models

A model is represented by a regular Python class that inherits from the Module class

# Forward Pass

It is the moment when the model makes predictions

```
# Sets learning rate - this is "eta" ~ the "n"-like
# Greek letter
lr = 0.1

# Step 0 - Initializes parameters "b" and "w" randomly
torch.manual_seed(42)
# Now we can create a model and send it at once to the
device
model = ManualLinearRegression().to(device)

# Defines a SGD optimizer to update the parameters
# (now retrieved directly from the model)
optimizer = optim.SGD(model.parameters(), lr=lr)

# Defines a MSE loss function
loss_fn = nn.MSELoss(reduction='mean')

# Defines number of epochs
n_epochs = 1000
```

```
for epoch in range(n_epochs):
    model.train() # What is this?!!

    # Step 1 - Computes model's predicted output - forward pass
    # No more manual prediction!
    yhat = model(x_train_tensor)

    # Step 2 - Computes the loss
    loss = loss_fn(yhat, y_train_tensor)

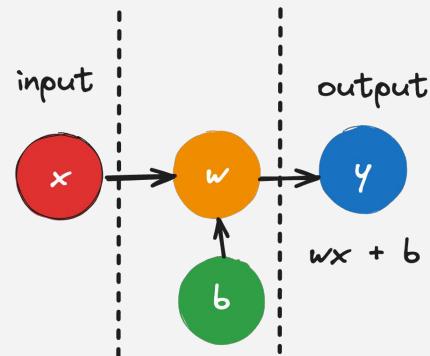
    # Step 3 - Computes gradients for both "b" and "w"parameters
    loss.backward()

    # Step 4 - Updates parameters using gradients and
    # the learning rate
    optimizer.step()
    optimizer.zero_grad()

    # We can also inspect its parameters using its state_dict
print(model.state_dict())
OrderedDict([('b', tensor([1.0235], device='cuda:0')),
            ('w', tensor([1.9690], device='cuda:0'))])
```

# Sequential Models

Our model was simple enough. You may be thinking: "Why even bother to build a class for it?!" Well, you have a point...



```
torch.manual_seed(42)
# Alternatively, you can use a Sequential model
model = nn.Sequential(nn.Linear(1, 1)).to(device)

model.state_dict()
OrderedDict([('0.weight', tensor([[0.7645]]), device='cuda:0')),
            ('0.bias', tensor([0.8300], device='cuda:0'))])
```

```
device = 'cuda' if torch.cuda.is_available() else 'cpu'

# Our data was in Numpy arrays, but we need to transform them
# into PyTorch's Tensors and then we send them to the
# chosen device
x_train_tensor = torch.as_tensor(x_train).float().to(device)
y_train_tensor = torch.as_tensor(y_train).float().to(device)
```

# Data Preparation

```
# Sets learning rate - this is "eta" ~ the "n"-like Greek letter
lr = 0.1

torch.manual_seed(42)
# Now we can create a model and send it at once to the device
model = nn.Sequential(nn.Linear(1, 1)).to(device)

# Defines a SGD optimizer to update the parameters
# (now retrieved directly from the model)
optimizer = optim.SGD(model.parameters(), lr=lr)

# Defines a MSE loss function
loss_fn = nn.MSELoss(reduction='mean')
```

# Model Configuration

```
# Defines number of epochs
n_epochs = 1000

for epoch in range(n_epochs):
    # Sets model to TRAIN mode
    model.train()

    # Step 1 - Computes model's predicted output - forward pass
    yhat = model(x_train_tensor)

    # Step 2 - Computes the loss
    loss = loss_fn(yhat, y_train_tensor)

    # Step 3 - Computes gradients for both "b" and "w" parameters
    loss.backward()

    # Step 4 - Updates parameters using gradients and
    # the learning rate
    optimizer.step()
    optimizer.zero_grad()

print(model.state_dict())
OrderedDict([('0.weight', tensor([[1.9690]], device='cuda:0')),
            ('0.bias', tensor([1.0235], device='cuda:0'))])
```

# Training

## Weeks 10 and 11 Fundamentals of Deep Learning

 PDF

- Outline  Video
- The perceptron  Video
- Building Neural Networks  Video
- Matrix Dimension  Video
- Applying Neural Networks  Video
- Training a Neural Networks  Video
- Backpropagation with Pencil & Paper  Video
- Learning rate & Batch Size  Video
- Exponentially Weighted Average  Video
- Adam, Momentum, RMSProp, Learning Rate Decay  Video
- Hands on 🔥
  - TensorFlow Crash Course  Notebook
  - Better Learning - Part I  Notebook



ppgeecmachinelearning

Public

Repository for EEC1509, a graduate course on PPgECEC about Machine Learning

Jupyter Notebook

41

21