

PPGEEC2318

Machine Learning

Rethinking the training loop: a simple classification problem

Part 01

Ivanovitch Silva

ivanovitch.silva@ufrn.br



- ▼ Part I: Fundamentals
 - › Chapter 0: Visualizing Gradient Descent
 - › Chapter 1: A Simple Regression Problem
 - › Chapter 2: Rethinking the Training Loop
 - › Chapter 2.1: Going Classy
 - › Chapter 3: A Simple Classification Problem
- › Part II: Computer Vision
- › Part III: Sequences
- › Part IV: Natural Language Processing

$$y = b + wx + \epsilon$$



Run

Data Generation

Data Preparation

Model Configuration

v0	v0	v0	v0	1
v0	v0	v1	v1	2
v0	v1	v1	v2	3
v0	v1	v1	v3	4
v0	v2	v2	v4	5

v0
Numpy
manual split

v0
x_train
y_train
as a tensor
(GPU)

v1
TensorDataset
DataLoader

v2
TensorDataset
DataLoader
Random Split

v0
sequential
optimizer
loss

v1
sequential
optimizer
loss
make_train_step_fn

v2
sequential
optimizer
loss
make_train_step_fn
make_val_step_fn

Model Training

Saving and Loading
Models

Deploy and Make
Predictions


v0
model.train()
yhat
loss
backward
optimizer update

v1
high-order function
build a training
step function
fn(x_train_tensor,
y_train_tensor)

v2
build a training
step function
mini-batch
dataloader

v3
build a training
step function
mini-batch function
dataloader

v4
build a training
step function
mini-batch function
dataloader
validation function



High-Order Functions

Higher-order functions in Python are functions that either **take other functions** as arguments or **return functions** as their results. This concept is key in functional programming and enhances flexibility and expressivity in your code.

Functions Accepting Functions as Arguments

```
def square(x):  
    return x * x
```

```
numbers = [1, 2, 3, 4, 5]  
squares = map(square, numbers)  
print(list(squares)) # Output: [1, 4, 9, 16, 25]
```

High-Order Functions

Functions Returning Functions

```
def multiplier(n):  
    def inner(x):  
        return x * n  
    return inner  
  
double = multiplier(2)  
triple = multiplier(3)  
  
print(double(5))    # Output: 10  
print(triple(5))    # Output: 15
```

Why Use High-Order Functions?

- **Code Reusability:** Using higher-order functions, you can write more generic and reusable code. For instance, you can write a general-purpose filter function that can be applied with different criterion functions.
- **Functional Programming:** Higher-order functions are a key component of functional programming, a paradigm emphasizing immutability and the use of functions to transform data.
- **Readability and Expressiveness:** They allow developers to express complex ideas more clearly and concisely.



Model lives
on GPU

Dataset

```
from torch.utils.data import Dataset

class CustomDataset(Dataset):
    def __init__(self, x_tensor, y_tensor):
        self.x = x_tensor
        self.y = y_tensor

    def __getitem__(self, index):
        return (self.x[index], self.y[index])

    def __len__(self):
        return len(self.x)

# Wait, is this a CPU tensor now? Why? Where is .to(device)?
x_train_tensor = torch.from_numpy(x_train).float()
y_train_tensor = torch.from_numpy(y_train).float()

train_data = CustomDataset(x_train_tensor, y_train_tensor)
print(train_data[0])
```




TensorDataset & DataLoader

```
# Our data was in Numpy arrays, but we need to transform them into
# PyTorch's Tensors
x_train_tensor = torch.from_numpy(x_train).float()
y_train_tensor = torch.from_numpy(y_train).float()

# Builds Dataset
train_data = TensorDataset(x_train_tensor, y_train_tensor)

# Builds DataLoader
train_loader = DataLoader(dataset=train_data,
                           batch_size=16,
                           shuffle=True)
```



TensorDataset & DataLoader

```
# Builds tensors from numpy arrays BEFORE split
x_tensor = torch.from_numpy(x).float()
y_tensor = torch.from_numpy(y).float()

# Builds dataset containing ALL data points
dataset = TensorDataset(x_tensor, y_tensor)

# Performs the split
ratio = .8
n_total = len(dataset)
n_train = int(n_total * ratio)
n_val = n_total - n_train

train_data, val_data = random_split(dataset, [n_train, n_val])

# Builds a loader of each set
train_loader = DataLoader(dataset=train_data,
                           batch_size=16, shuffle=True)
val_loader = DataLoader(dataset=val_data, batch_size=16)
```



Model Configuration

```
device = 'cuda' if torch.cuda.is_available() else 'cpu'

# Sets learning rate - this is "eta" ~ the "n" like Greek letter
lr = 0.1

torch.manual_seed(42)
# Now we can create a model and send it at once to the device
model = nn.Sequential(nn.Linear(1, 1)).to(device)

# Defines a SGD optimizer to update the parameters (now retrieved directly from
the model)
optimizer = optim.SGD(model.parameters(), lr=lr)

# Defines a MSE loss function
loss_fn = nn.MSELoss(reduction='mean')

# Creates the train_step function for our model, loss function and optimizer
train_step_fn = make_train_step_fn(model, loss_fn, optimizer)

# Creates the val_step function for our model and loss function
val_step_fn = make_val_step_fn(model, loss_fn)
```


Model Configuration

```
def make_train_step_fn(model, loss_fn, optimizer):  
    # Builds function that performs a step in the train loop  
    def perform_train_step_fn(x, y):  
        # Sets model to TRAIN mode  
        model.train()  
  
        # Step 1 - Computes our model's predicted output - forward pass  
        yhat = model(x)  
        # Step 2 - Computes the loss  
        loss = loss_fn(yhat, y)  
        # Step 3 - Computes gradients for both "a" and "b" parameters  
        loss.backward()  
        # Step 4 - Updates parameters using gradients and the learning rate  
        optimizer.step()  
        optimizer.zero_grad()  
  
        # Returns the loss  
        return loss.item()  
  
    # Returns the function that will be called inside the train loop  
    return perform_train_step_fn
```

Model Configuration

```
def make_val_step_fn(model, loss_fn):  
    # Builds function that performs a step in the validation loop  
    def perform_val_step_fn(x, y):  
        # Sets model to EVAL mode  
        model.eval()  
  
        # Step 1 - Computes our model's predicted output - forward pass  
        yhat = model(x)  
        # Step 2 - Computes the loss  
        loss = loss_fn(yhat, y)  
        # There is no need to compute Steps 3 and 4, since we don't update  
        # parameters during evaluation  
        return loss.item()  
  
    return perform_val_step_fn
```

Training

```
# Defines number of epochs
n_epochs = 200


losses = []
val_losses = []


for epoch in range(n_epochs):
    # inner loop
    loss = mini_batch(device, train_loader, train_step_fn)
    losses.append(loss)

    # VALIDATION
    # no gradients in validation!
    with torch.no_grad():
        val_loss = mini_batch(device, val_loader, val_step_fn)
        val_losses.append(val_loss)
```


Training

```
def mini_batch(device, data_loader, step_fn):  
    mini_batch_losses = []  
    for x_batch, y_batch in data_loader:  
        x_batch = x_batch.to(device)  
        y_batch = y_batch.to(device)  
  
        mini_batch_loss = step_fn(x_batch, y_batch)  
        mini_batch_losses.append(mini_batch_loss)  
  
    loss = np.mean(mini_batch_losses)  
    return loss
```





LOADING...

Saving and Loading Models

```
checkpoint = {'epoch': n_epochs,  
              'model_state_dict': model.state_dict(),  
              'optimizer_state_dict': optimizer.state_dict(),  
              'loss': losses,  
              'val_loss': val_losses}  
  
torch.save(checkpoint, 'model_checkpoint.pth')
```

```
checkpoint = torch.load('model_checkpoint.pth', weights_only=False)  
  
model.load_state_dict(checkpoint['model_state_dict'])  
optimizer.load_state_dict(checkpoint['optimizer_state_dict'])  
  
saved_epoch = checkpoint['epoch']  
saved_losses = checkpoint['loss']  
saved_val_losses = checkpoint['val_loss']  
  
model.train() # always use TRAIN for resuming training
```


Make Predictions

```
checkpoint = torch.load('model_checkpoint.pth', weights_only=False)

model.load_state_dict(checkpoint['model_state_dict'])

new_inputs = torch.tensor([[.20], [.34], [.57]])

model.eval() # always use EVAL for fully trained models!
model(new_inputs.to(device))

tensor([[1.4174],
        [1.6896],
        [2.1366]], device='cuda:0')
```



Going Classy

Object-Oriented Programming – OOP

week04a.ipynb, week04b.ipynb

```
139 title="Home"
140 target="_blank"
141 rule="noopener noreferrer"
142 href={trackUrl}
143 >
144 </a>
145 </li>
146 </ul>
147 </div>
148 </div>
149 </div>
150 </div>
151 </div>
152 </div>
153 </div>
154 </div>
155 </div>
156 </div>
157 </div>
158 </div>
159 </div>
160 </div>
161 </div>
162 </div>
163 </div>
164 </div>
165 </div>
166 </div>
167 </div>
168 </div>
169 </div>
170 </div>
171 </div>
172 </div>
173 </div>
174 </div>
175 </div>
176 </div>
177 </div>
178 </div>
179 </div>
180 </div>
181 </div>
182 </div>
183 </div>
184 </div>
185 </div>
186 </div>
187 </div>
188 </div>
189 </div>
190 </div>
191 </div>
192 </div>
193 </div>
194 </div>
195 </div>
196 </div>
197 </div>
198 </div>
199 </div>
200 </div>
201 </div>
202 </div>
203 </div>
204 </div>
205 </div>
206 </div>
207 </div>
208 </div>
209 </div>
```