

PPGEEC2318

Machine Learning

Rock, Paper, Scissors

Ivanovitch Silva

ivanovitch.silva@ufrn.br



Daniel Voigt Godoy

Deep Learning with PyTorch Step-by-Step



A Beginner's Guide



Chapter 6: Rock, Paper, Scissors

Spoilers

- > Jupyter Notebook
- > Rock, Paper, Scissors...
- > Data Preparation

Three-Channel Convolutions

Fancier Model

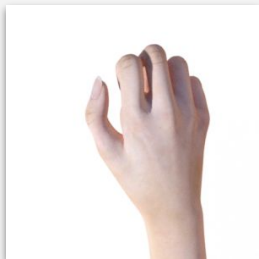
- > Dropout
- > Model Configuration
- > Model Training
- > Learning Rates

Putting It All Together

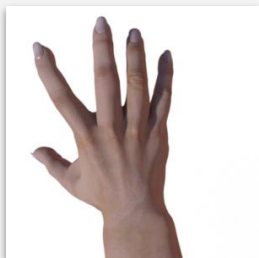
Recap

Agenda

1. **Standardize** an image dataset
2. **train** a model to predict **rock, paper, scissors** poses from hand images
3. use **dropout** layers to **regularize** the model



Rock

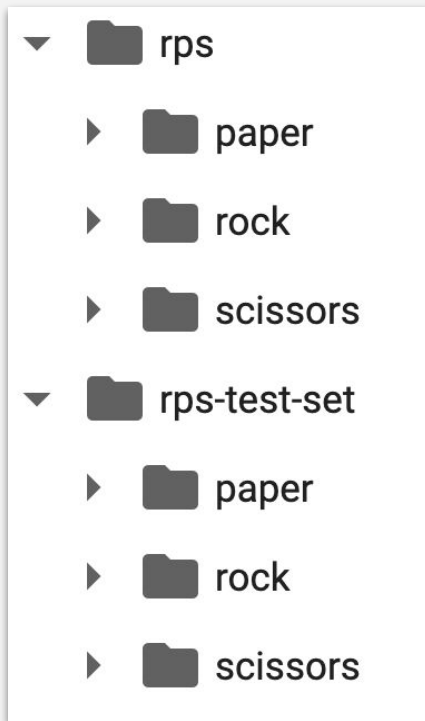
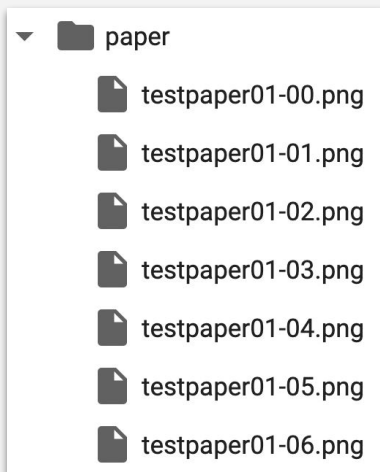
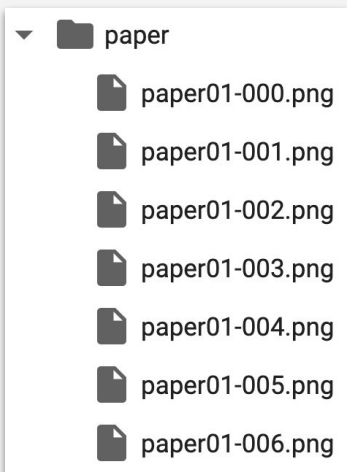


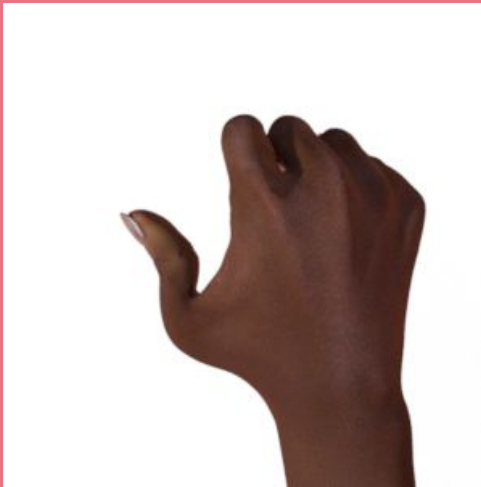
Paper



Scissors

The dataset contains **2,892 images (2,520 train, 372 test)** of diverse hands in the typical rock, paper, and scissors poses against a white background. This is a synthetic dataset as well since the images were generated using CGI techniques. Each image is **300x300 pixels** in size and has **four channels (RGBA)**.





If the images are colored

We need to standardize the three channels (RGB)

Find the $\langle \text{mean}, \text{std} \rangle$ for each channel

and to limit them to $\langle 0, 1 \rangle$

Only for train dataset!!! Avoid data leakage!!!

Data Preparation

ImageFolder

```
# Compose a sequence of preprocessing transforms
# 1) Resize images to 28x28 pixels
# 2) Ensure output is a PIL/torchvision Image (dropping any alpha channel)
# 3) Convert pixel values to float32 and scale from [0-255] to [0.0-1.0]
temp_transform = Compose([
    Resize(28),                    # Resize each image to 28x28
    ToImage(),                     # Convert tensor back to PIL Image (enforces RGB)
    ToDtype(torch.float32, scale=True) # Cast to float32 and normalize pixel range
])

# Create an ImageFolder dataset from the 'rps' directory
# Images are grouped by subfolder name as class labels, and each image is transformed
temp_dataset = ImageFolder(
    root='rps',
    transform=temp_transform      # Apply the preprocessing pipeline to every image
)
```


Data Preparation

ImageFolder



```
# Get total number of samples in the dataset
dataset_size = len(temp_dataset)
print(f"Dataset size: {dataset_size} images")
```

```
# Get number of classes
num_classes = len(temp_dataset.classes)
print(f"Number of classes: {num_classes}")
```

```
Dataset size: 2520 images
Number of classes: 3
```



```
# the second element of this tuple is the label
temp_dataset[0][0].shape, temp_dataset[0][1]
```

```
(torch.Size([3, 28, 28]), 0)
```

```
temp_dataset[0][0]
```

```
Image([[ [1.0000, 1.0000, 1.0000, ..., 1.0000, 1.0000, 1.0000],
        [1.0000, 1.0000, 1.0000, ..., 1.0000, 1.0000, 1.0000],
        [1.0000, 1.0000, 1.0000, ..., 1.0000, 1.0000, 1.0000],
        ...,
        [1.0000, 1.0000, 1.0000, ..., 0.9961, 0.9961, 0.9961],
        [1.0000, 1.0000, 1.0000, ..., 0.9961, 0.9961, 0.9922],
        [1.0000, 1.0000, 1.0000, ..., 0.9961, 0.9922, 0.9922]],
       [[ [1.0000, 1.0000, 1.0000, ..., 1.0000, 1.0000, 1.0000],
        [1.0000, 1.0000, 1.0000, ..., 1.0000, 1.0000, 1.0000],
        [1.0000, 1.0000, 1.0000, ..., 1.0000, 1.0000, 1.0000],
        ...,
        [1.0000, 1.0000, 1.0000, ..., 0.9961, 0.9961, 0.9961],
        [1.0000, 1.0000, 1.0000, ..., 0.9961, 0.9961, 0.9922],
        [1.0000, 1.0000, 1.0000, ..., 0.9961, 0.9922, 0.9922]],
       [[ [1.0000, 1.0000, 1.0000, ..., 1.0000, 1.0000, 1.0000],
        [1.0000, 1.0000, 1.0000, ..., 1.0000, 1.0000, 1.0000],
        [1.0000, 1.0000, 1.0000, ..., 1.0000, 1.0000, 1.0000],
        ...,
        [1.0000, 1.0000, 1.0000, ..., 0.9961, 0.9961, 0.9961],
        [1.0000, 1.0000, 1.0000, ..., 0.9961, 0.9961, 0.9922],
        [1.0000, 1.0000, 1.0000, ..., 0.9961, 0.9922, 0.9922]]], )
```

Data Preparation

Standardization

```
temp_loader = DataLoader(temp_dataset, batch_size=16)
# Each column represents a channel
# first row is the number of data points
# second row is the the sum of mean values
# third row is the sum of standard deviations
first_images, first_labels = next(iter(temp_loader))
Architecture.statistics_per_channel(first_images, first_labels)

tensor([[16.0000, 16.0000, 16.0000],
        [13.8748, 13.3048, 13.1962],
        [ 3.0507,  3.8268,  3.9754]])

# We can leverage the loader_apply() method to get the sums for the whole dataset:
results = Architecture.loader_apply(temp_loader,
                                    Architecture.statistics_per_channel)

tensor([[2520.0000, 2520.0000, 2520.0000],
        [2142.5356, 2070.0806, 2045.1444],
        [ 526.3025,  633.0677,  669.9556]])
```


Data Preparation

Standardization

```
temp_loader = DataLoader(temp_dataset, batch_size=16)

# we can compute the average mean value and the average standard deviation, per channel.
# Better yet, let's make it a method that takes a data loader and
# returns an instance of the Normalize() transform
normalizer = Architecture.make_normalizer(temp_loader)
normalizer

Normalize(mean=tensor([0.8502, 0.8215, 0.8116]),
          std=tensor([0.2089, 0.2512, 0.2659]))
```

Data Preparation

The real dataset

```
composer = Compose([
    Resize(28),           # Resize to 28x28
    ToImage(),            # Convert to PIL Image in RGB
    ToDtype(torch.float32, scale=True), # Cast to float32 and normalize to [0,1]
    normalizer            # Apply custom normalization transform
])

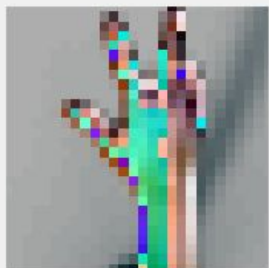
# Instantiate training and validation datasets from folders:
# - 'rps' contains subfolders per class for training
# - 'rps-test-set' likewise for validation
train_data = ImageFolder(root='rps', transform=composer)
val_data   = ImageFolder(root='rps-test-set', transform=composer)

# Wrap datasets in DataLoaders for batching and shuffling:
# - batch_size=16 yields mini-batches of 16 images
# - shuffle=True randomizes training order each epoch
train_loader = DataLoader(train_data, batch_size=16, shuffle=True)
val_loader   = DataLoader(val_data,   batch_size=16) # no shuffle for validation
```

Scissors



Scissors



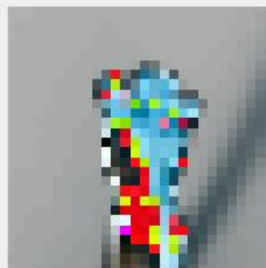
Paper



Paper



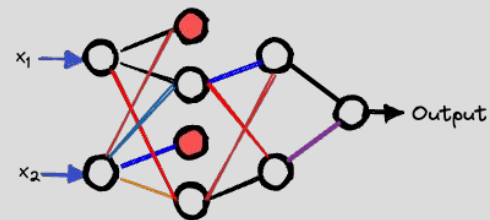
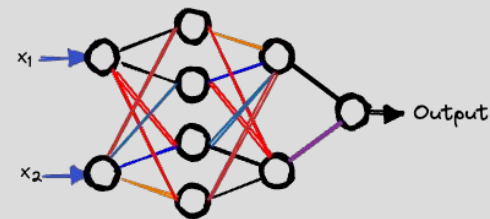
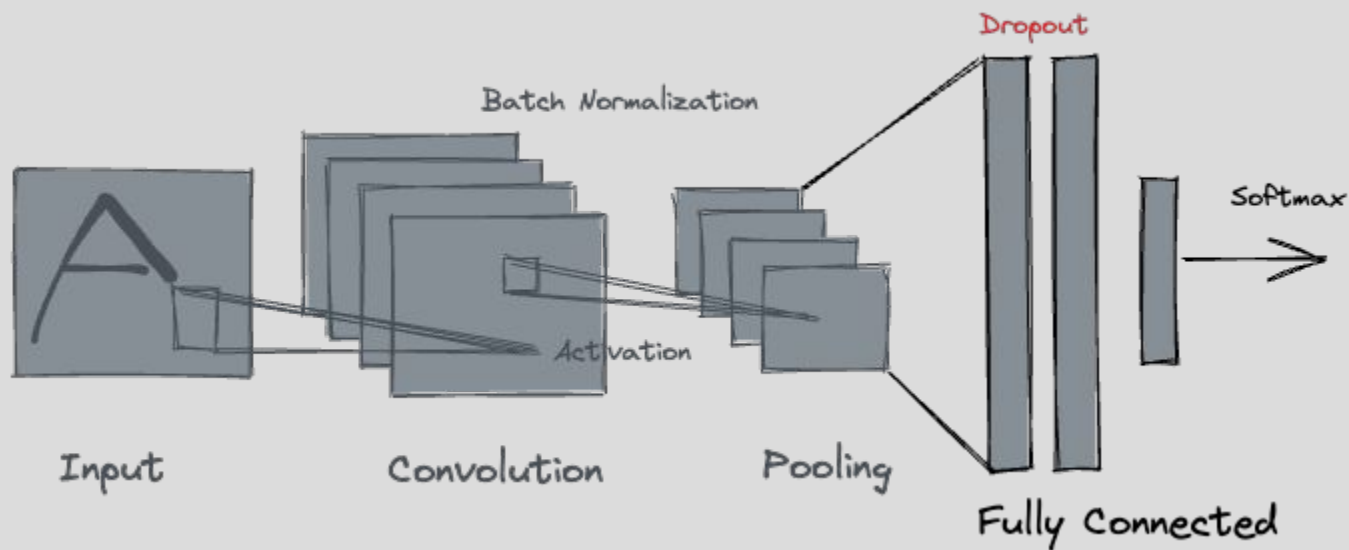
Rock



Rock



Dropout



It is a form of regularization

Reduces overfitting

Increases test/validation accuracy (sometimes at expense of training accuracy)

Randomly disconnects node from current layers to next layer with probability, p

Dropout (what's going on here?)



```
dropping_model = nn.Sequential(nn.Dropout(p=0.5))
```

```
spaced_points = torch.linspace(.1, 1.1, 11)
```

```
spaced_points
```

```
tensor([0.1000, 0.2000, 0.3000, 0.4000, 0.5000, 0.6000, 0.7000, 0.8000, 0.9000,  
        1.0000, 1.1000])
```

```
dropping_model.train()
```

```
output_train = dropping_model(spaced_points)
```

```
output_train
```

```
tensor([0.0000, 0.4000, 0.0000, 0.8000, 0.0000, 1.2000, 1.4000, 1.6000, 1.8000,  
        0.0000, 2.2000])
```

Dropout (what's going on here?)



```
F.linear(output_train, weight=torch.ones(11), bias=torch.tensor(0))  
tensor(9.4000)
```

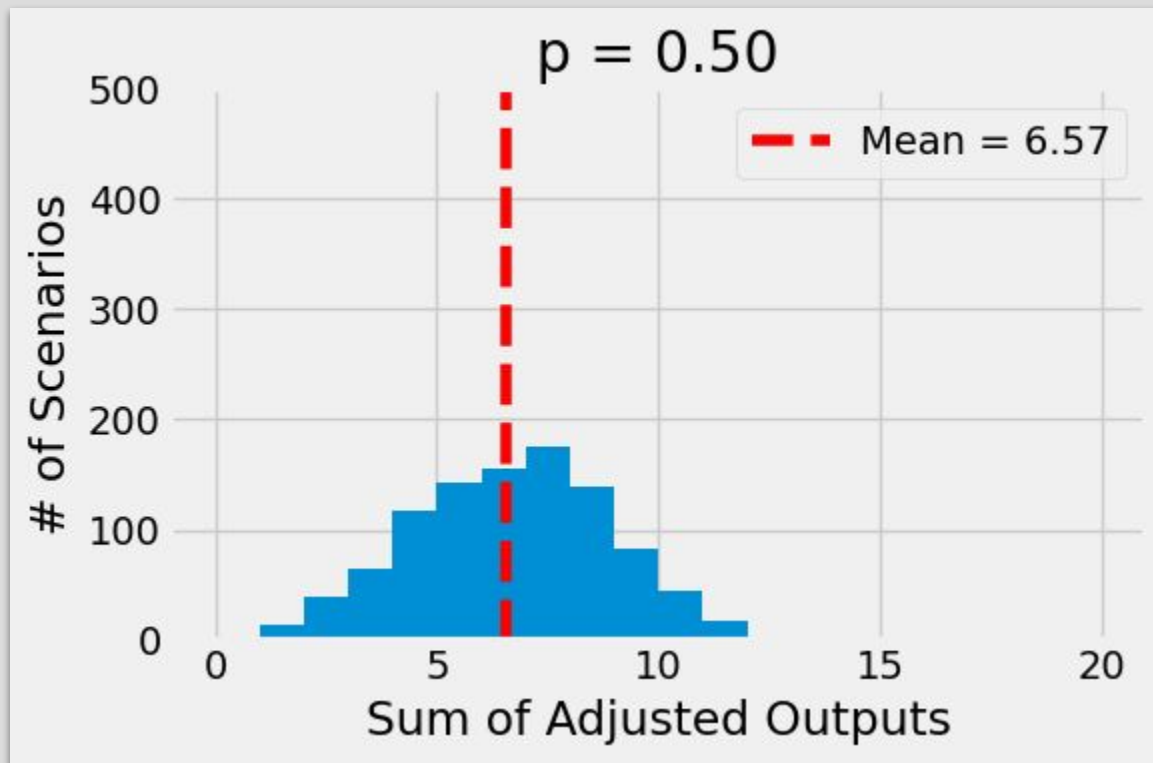
```
dropping_model.eval()  
output_eval = dropping_model(spaced_points)  
output_eval
```

```
tensor([0.1000, 0.2000, 0.3000, 0.4000, 0.5000, 0.6000, 0.7000, 0.8000, 0.9000,  
        1.0000, 1.1000])
```

```
F.linear(output_eval, weight=torch.ones(11), bias=torch.tensor(0))  
tensor(6.6000)
```

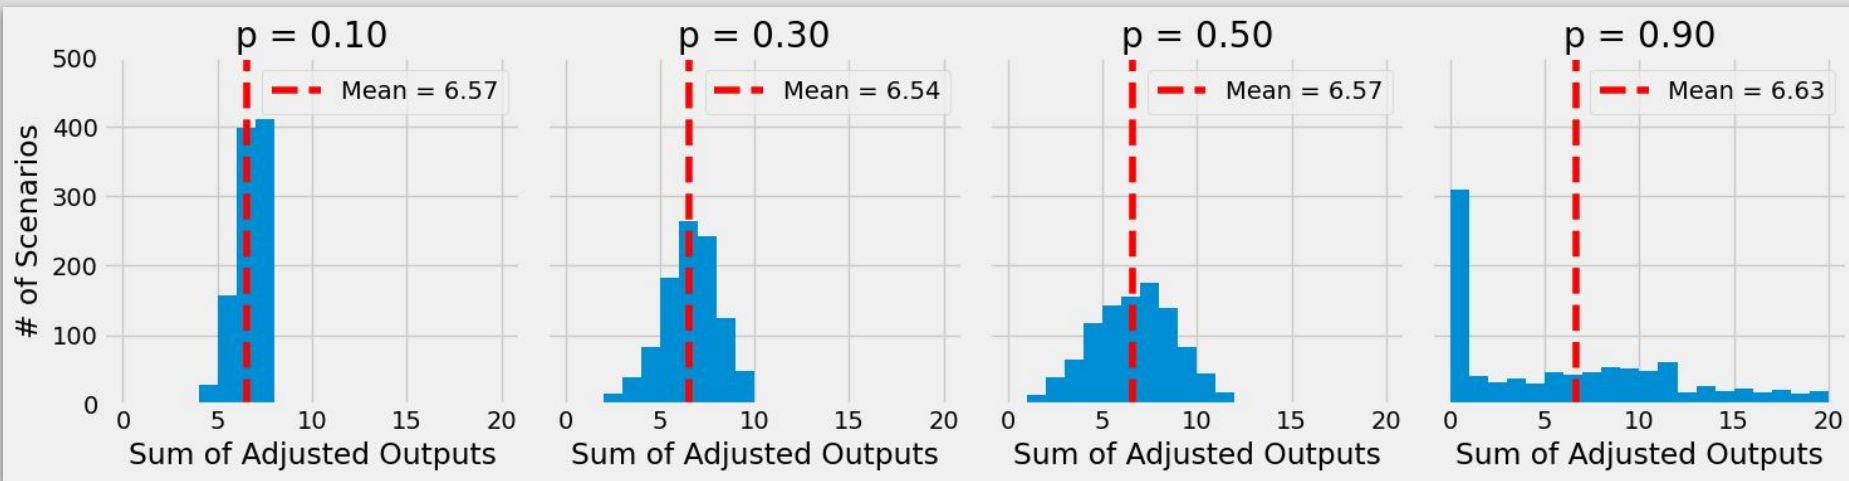
Dropout

Distribution of 1000 outputs



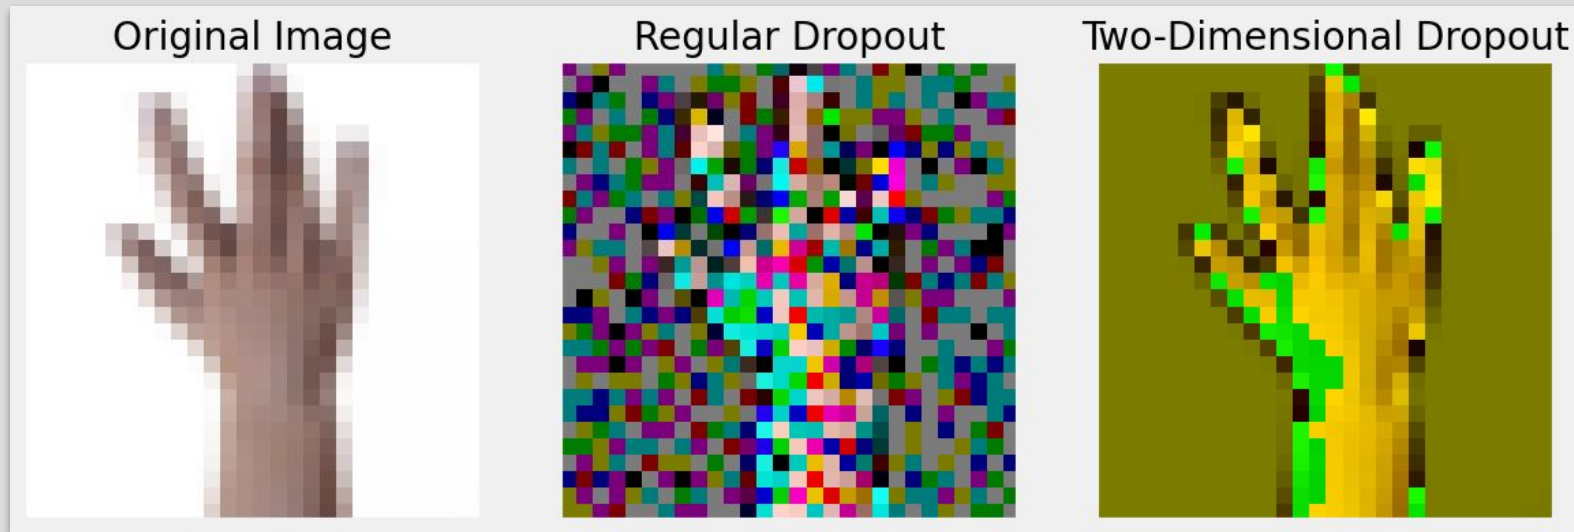
Dropout

Distribution of 1000 outputs



- For more typical dropout probabilities (like 30% or 50%), the distribution may take some more extreme values when compared to 10%
- The variance of the distribution of outputs grows with the dropout probability.
- A higher dropout probability makes it harder for your model to learn—that's what regularization does

Two dimensional dropout

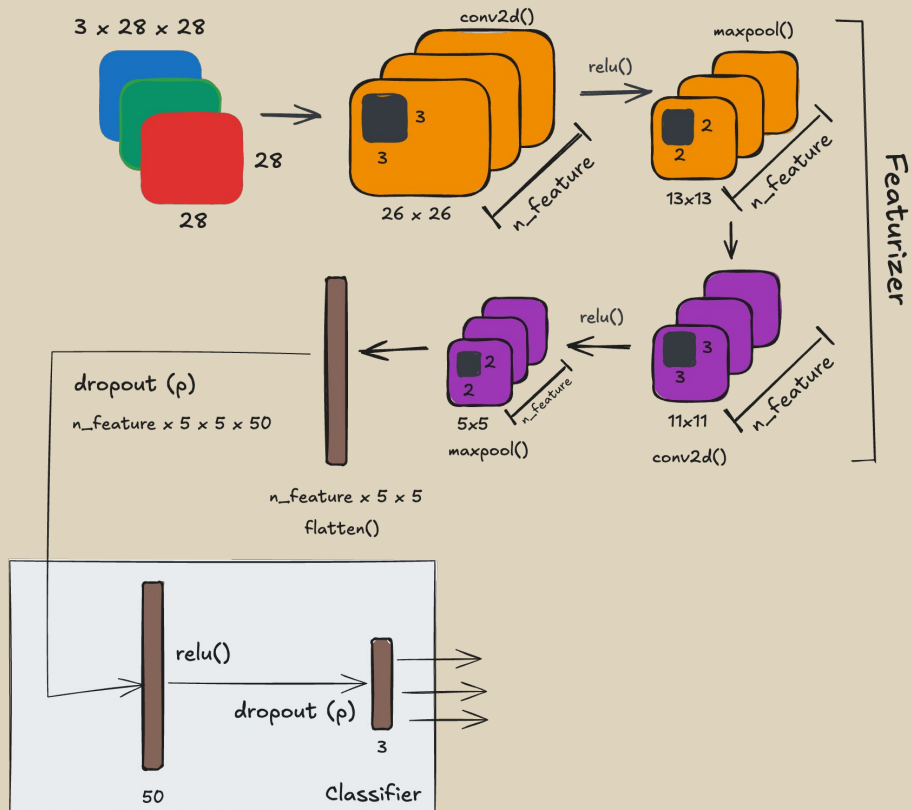


- It drops entire channels / filters.
- If a convolutional layer produces ten filters, a two-dimensional dropout with a probability of 50% would drop five filters (on average)
- The remaining filters would have all their pixel values left untouched.

Fancier Model

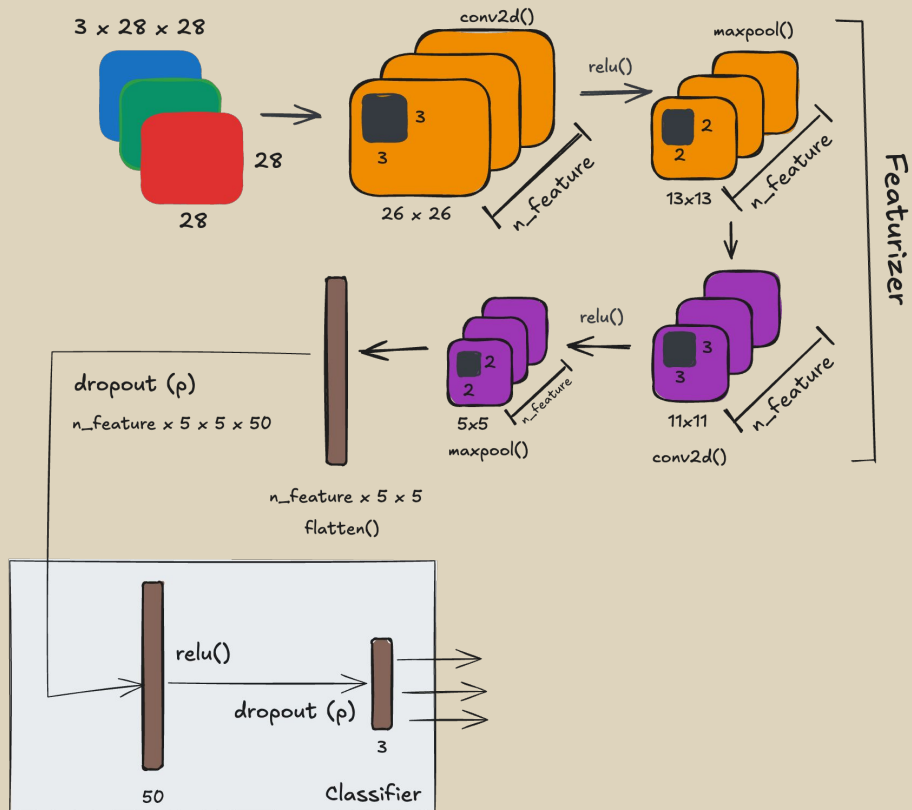
```
class CNN2(nn.Module):
    def __init__(self, n_feature, p=0.0):
        super(CNN2, self).__init__()
        self.n_feature = n_feature
        self.p = p
        # Creates the convolution layers
        self.conv1 = nn.Conv2d(in_channels=3,
                                out_channels=n_feature,
                                kernel_size=3)
        self.conv2 = nn.Conv2d(in_channels=n_feature,
                                out_channels=n_feature,
                                kernel_size=3)

        # Creates the linear layers
        # Where do this 5 * 5 come from?! Check it below
        self.fc1 = nn.Linear(n_feature * 5 * 5, 50)
        self.fc2 = nn.Linear(50, 3)
        # Creates dropout layers
        self.drop = nn.Dropout(self.p)
```



Fancier Model

```
def featurizer(self, x):
    # Featurizer
    # First convolutional block
    # 3@28x28 -> n_feature@26x26 -> n_feature@13x13
    x = self.conv1(x)
    x = F.relu(x)
    x = F.max_pool2d(x, kernel_size=2)
    # Second convolutional block
    # n_feature * @13x13 -> n_feature@11x11 -> n_feature@5x5
    x = self.conv2(x)
    x = F.relu(x)
    x = F.max_pool2d(x, kernel_size=2)
    # Input dimension (n_feature@5x5)
    # Output dimension (n_feature * 5 * 5)
    x = nn.Flatten()(x)
    return x
```

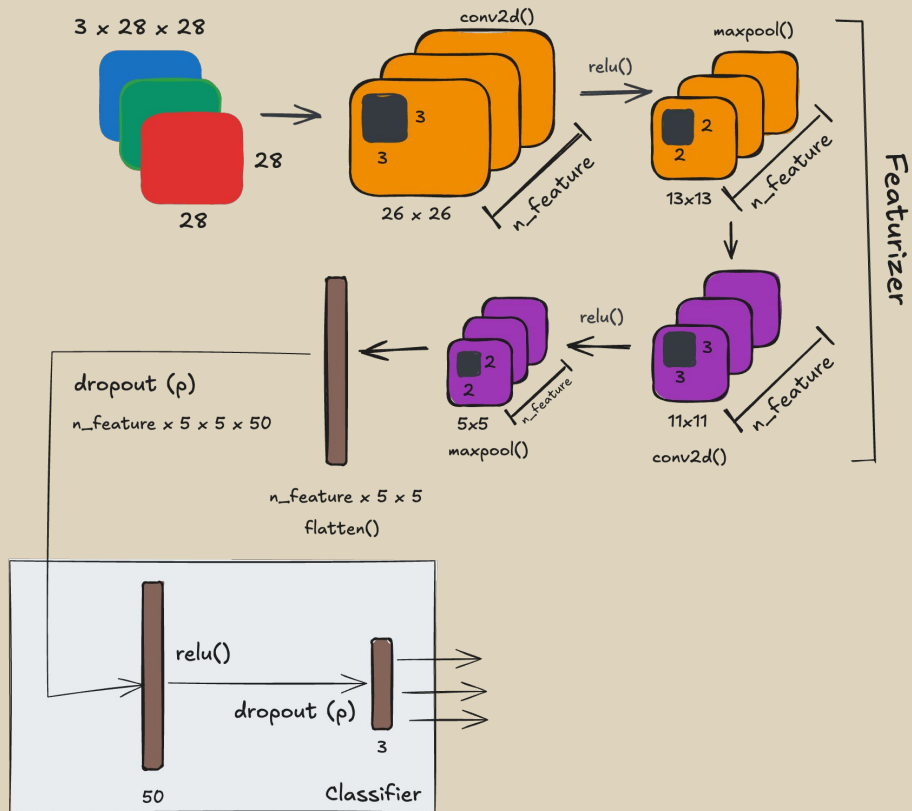


Fancier Model

```
def classifier(self, x):  
    # Classifier  
    # Hidden Layer  
    # Input dimension (n_feature * 5 * 5)  
    # Output dimension (50)  
    if self.p > 0:  
        x = self.drop(x)  
    x = self.fc1(x)  
    x = F.relu(x)  
    # Output Layer  
    # Input dimension (50)  
    # Output dimension (3)  
    if self.p > 0:  
        x = self.drop(x)  
    x = self.fc2(x)  
    return x
```

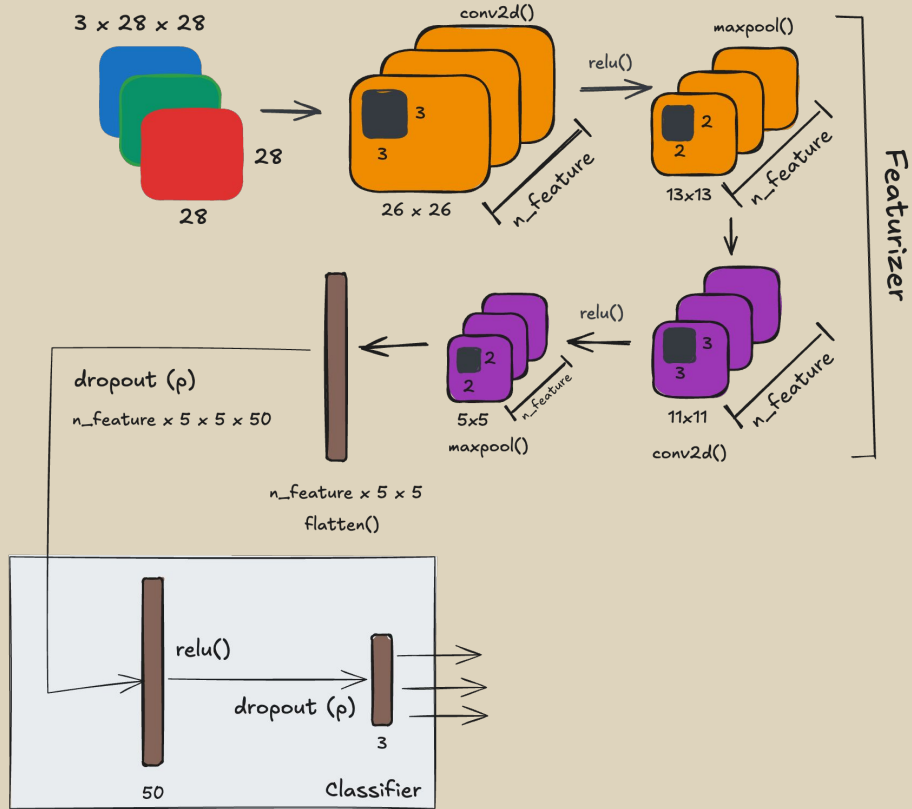


```
def forward(self, x):  
    x = self.featurizer(x)  
    x = self.classifier(x)  
    return x
```



Case Study

Model Configuration
Model Training
Accuracy
Regularizing Effect
Visualizing Filters



Model Config.

```
torch.manual_seed(13)

# Model/Architecture
model_cnn2 = CNN2(n_feature=5, p=0.3)

# Loss function
multi_loss_fn = nn.CrossEntropyLoss(reduction='mean')

# Optimizer
optimizer_cnn2 = optim.Adam(model_cnn2.parameters(), lr=3e-4)
```



Andrej Karpathy ✓
@karpathy

3e-4 is the best learning rate for Adam, hands down.

12:01 AM · Nov 24, 2016



31



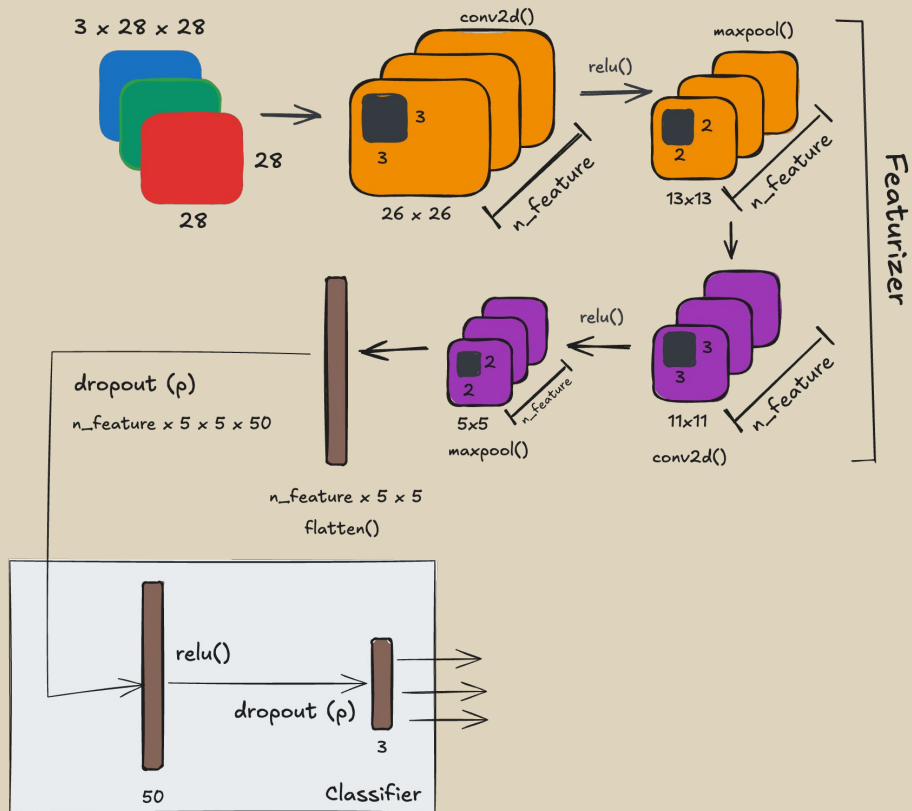
183



683



43



Model Train

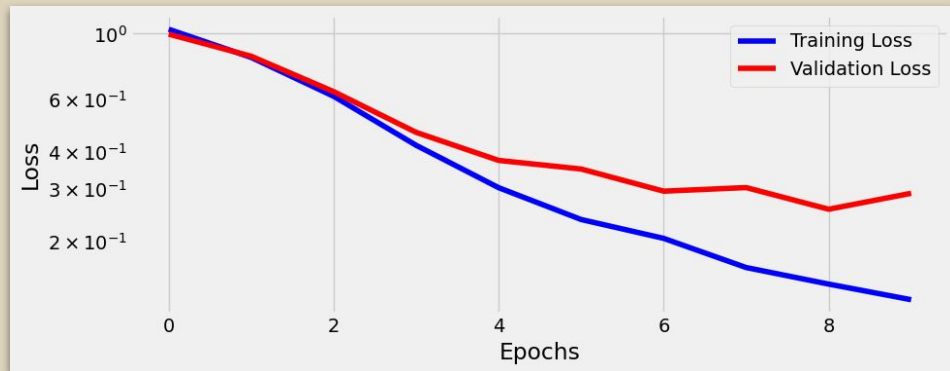


```
arch_cnn2 = Architecture(model_cnn2,  
                          multi_loss_fn,  
                          optimizer_cnn2)  
arch_cnn2.set_loaders(train_loader, val_loader)  
arch_cnn2.train(10)
```

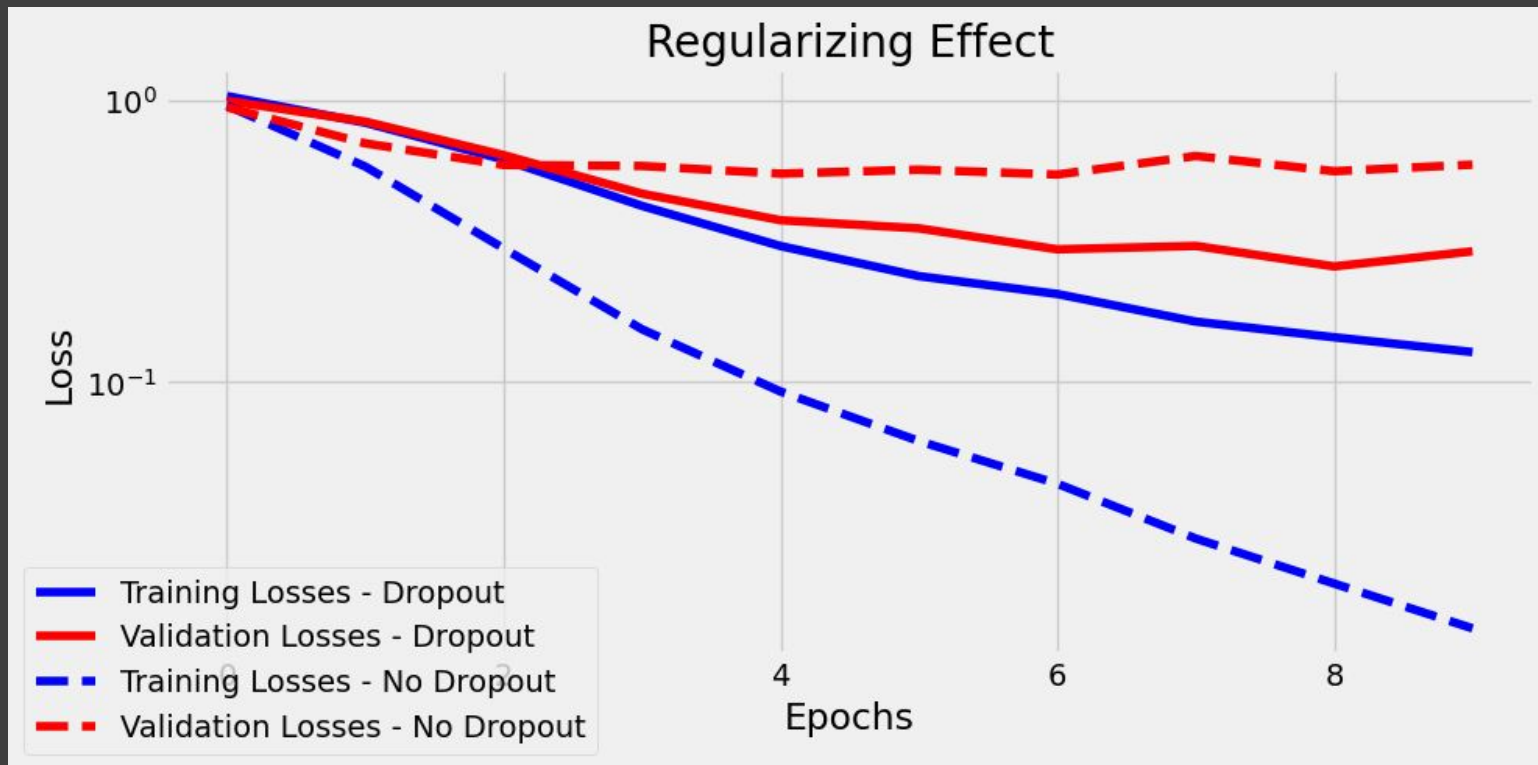
```
arch_cnn2.count_parameters()  
6823
```

```
Architecture.loader_apply(val_loader,  
                          arch_cnn2.correct)
```

```
tensor([[ 89, 124],  
        [118, 124],  
        [117, 124]])    87%
```



Regularizing Effect



Regularizing Effect



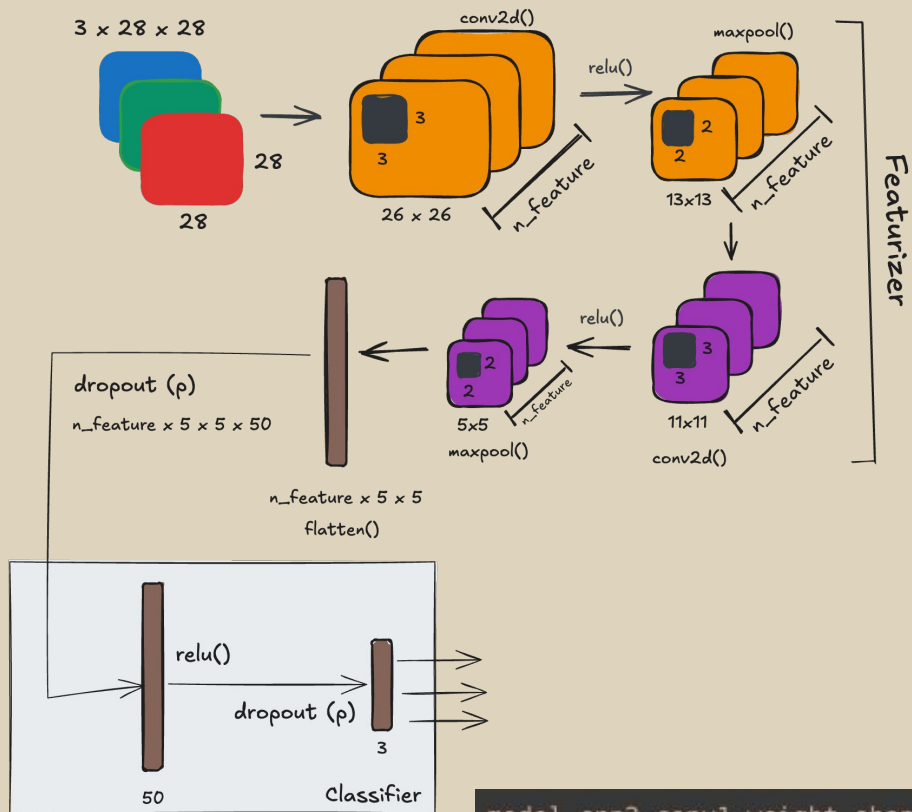
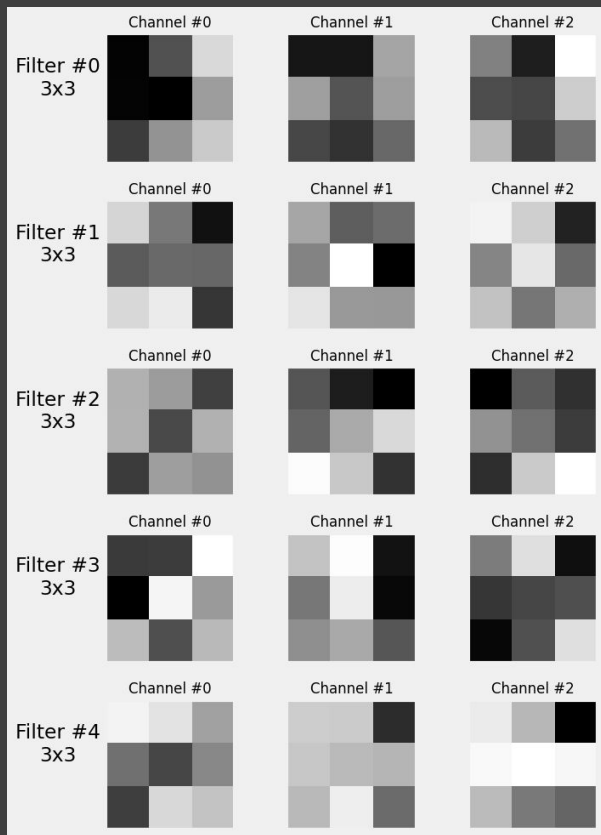
```
# no dropout
print(
    Architecture.loader_apply(train_loader, arch_cnn2_nodrop.correct).sum(axis=0),
    Architecture.loader_apply(val_loader, arch_cnn2_nodrop.correct).sum(axis=0)
)

tensor([2520, 2520]) tensor([292, 372])
1.0 0.7849462365591398

# with dropout
print(
    Architecture.loader_apply(train_loader, arch_cnn2.correct).sum(axis=0),
    Architecture.loader_apply(val_loader, arch_cnn2.correct).sum(axis=0)
)

tensor([2504, 2520]) tensor([326, 372])
0.9936507936507937 0.8763440860215054
```

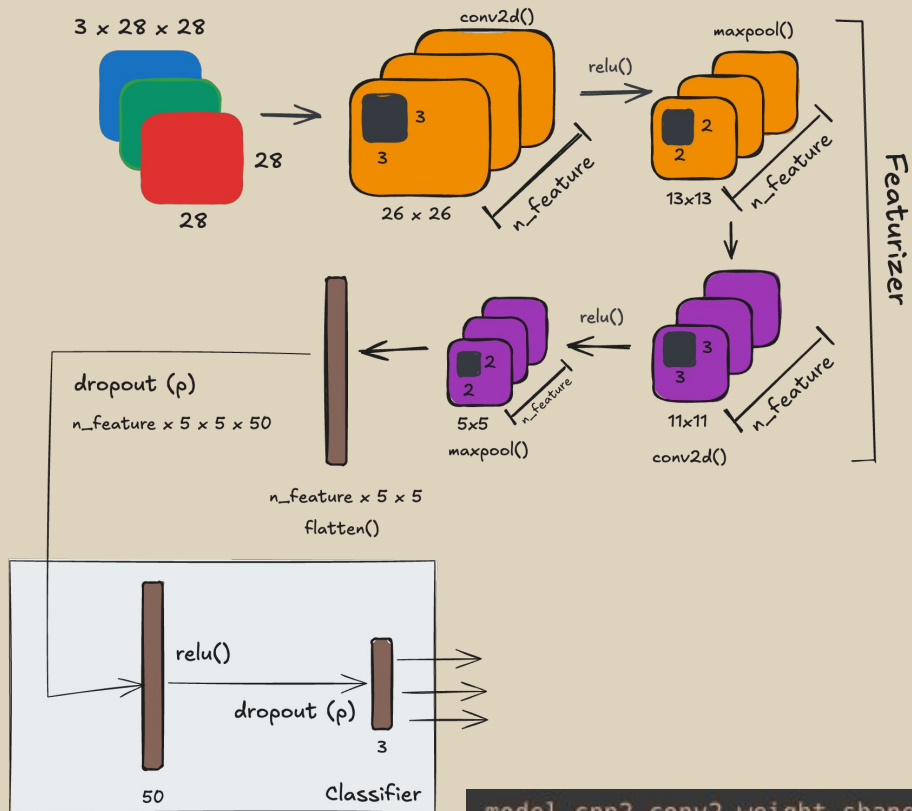
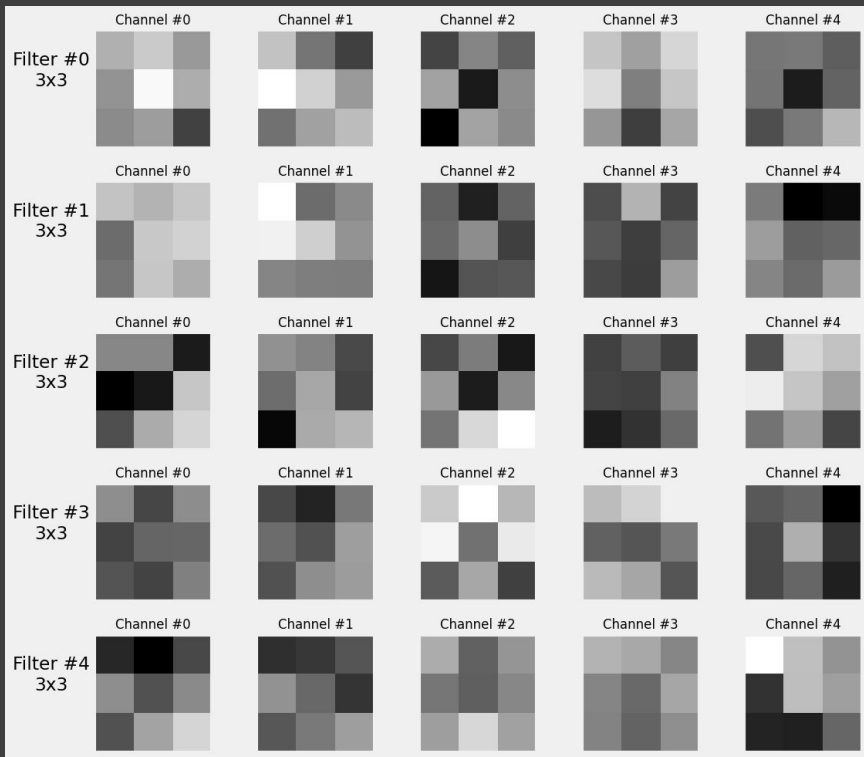
Visual. Filters



```
model_cnn2.conv1.weight.shape
```

```
torch.Size([5, 3, 3, 3])
```

Visual. Filters



```
model_cnn2.conv2.weight.shape
```

```
torch.Size([5, 5, 3, 3])
```

Cont.